

Overview of Machine Learning



AI, Machine Learning & Deep Learning

- **Artificial intelligence**

Computer technologies to simulate **human intelligence**.

- **Machine Learning**

[broader]

Subset of AI techniques used to make predictions by **learning from data** and **improving from experience**.

Learning process -> **Finding functions**.

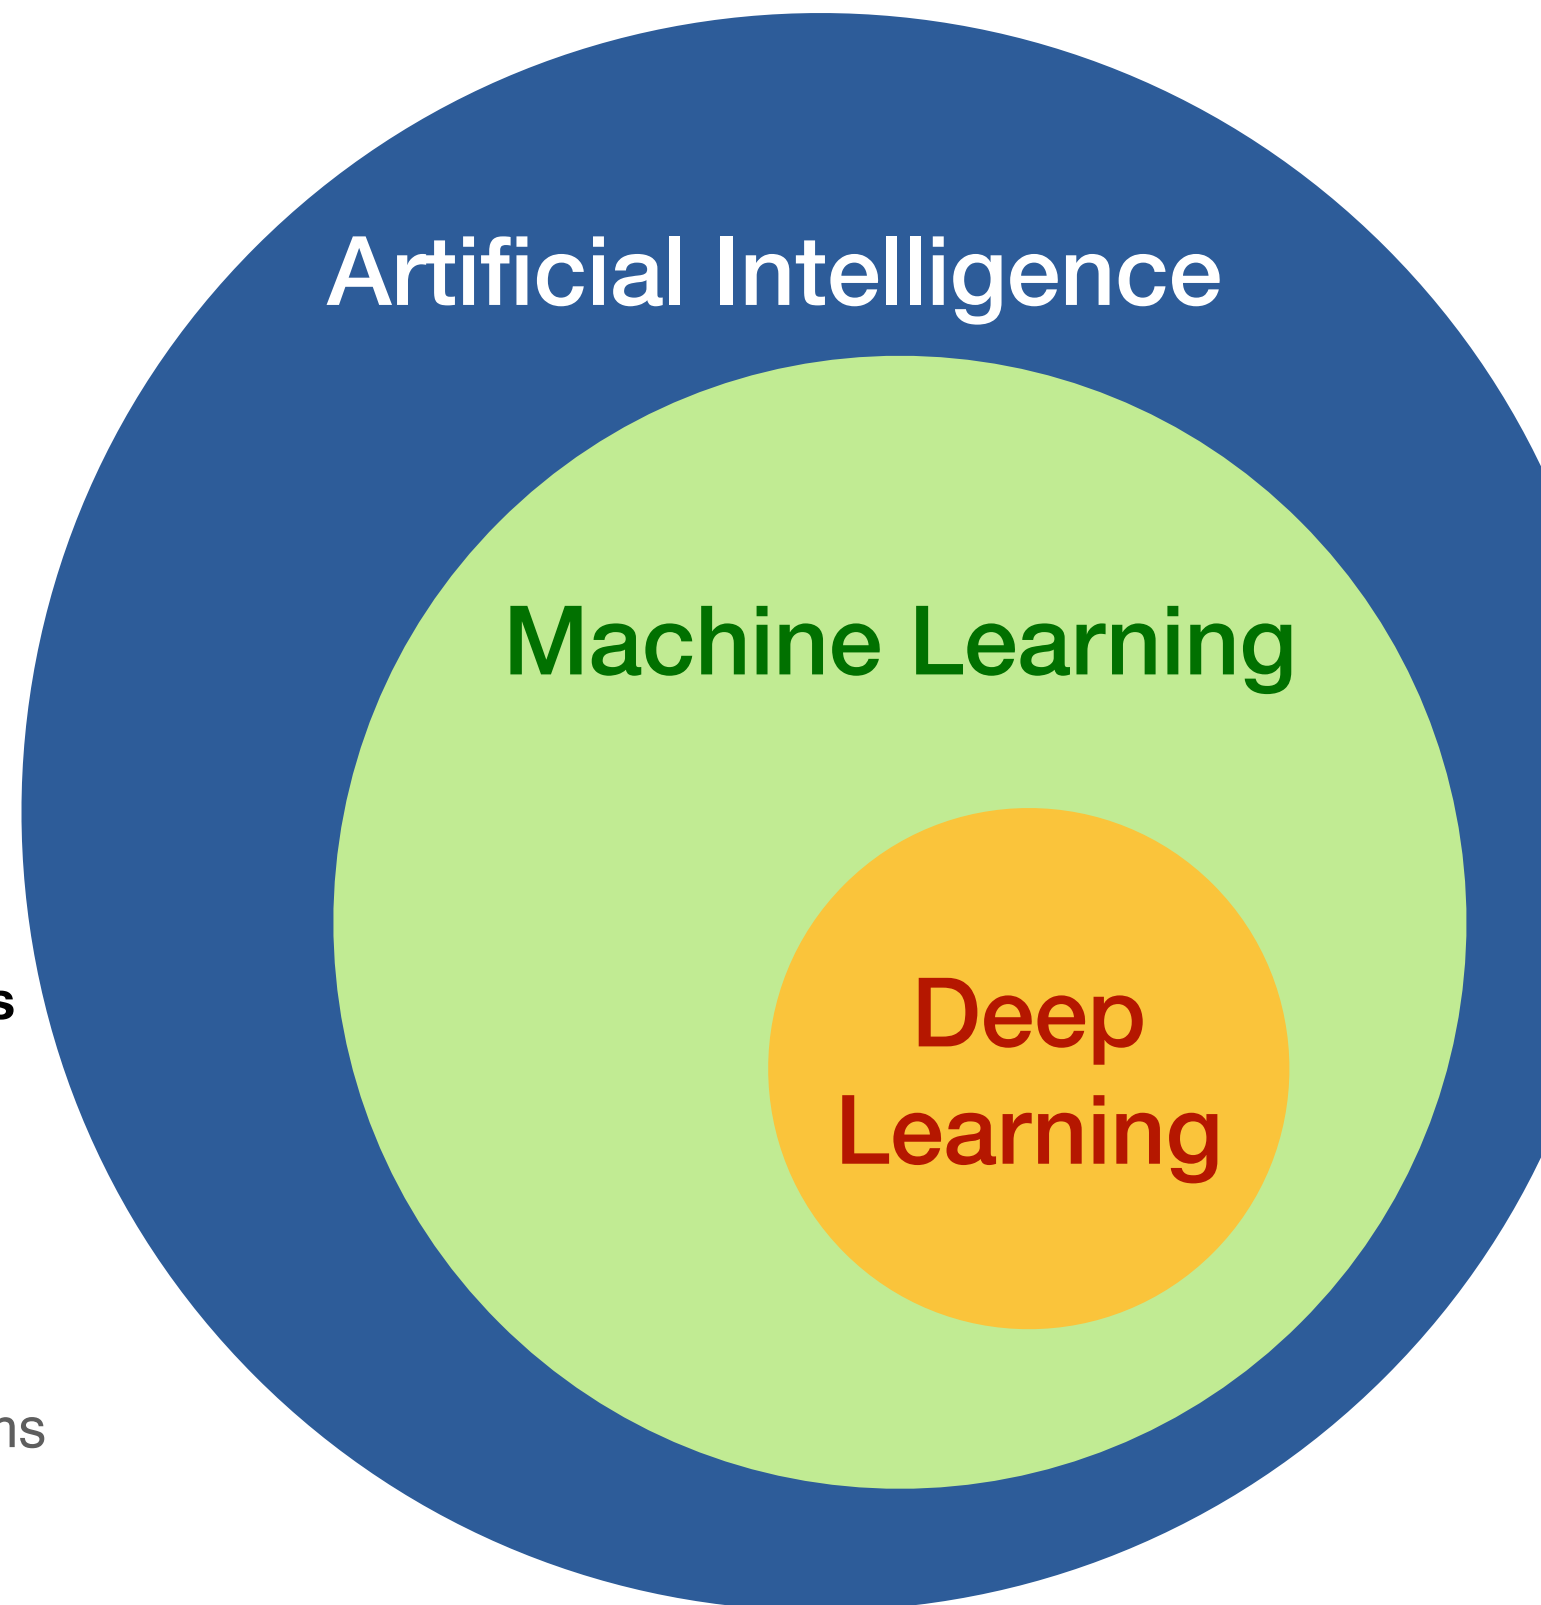
[narrower]

Classical **statistical learning algorithms**
(regression, clustering, SVM, random forests ...)

- **Deep Learning**

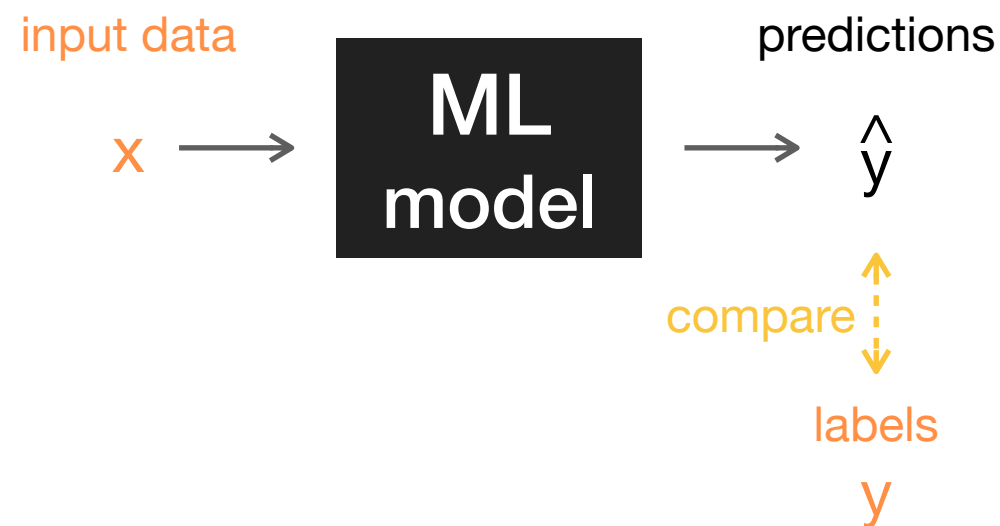
Subset of ML.

Neural network based learning algorithms



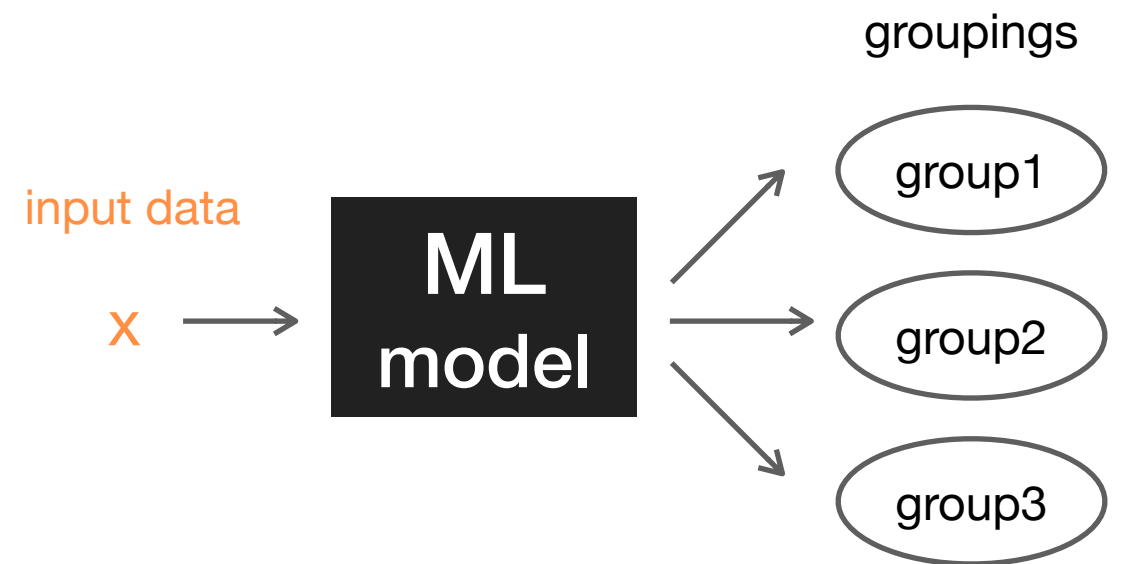
Types of Machine Learning

Supervised Learning



- Learning a function that maps an input to an output based on **labeled data**
- Goal : **make predictions** with high **accuracy**
- Applications : regression, classification
- Example algorithms : Linear / Logistic regression
Image classification networks
- Challenges : labeled data is expensive

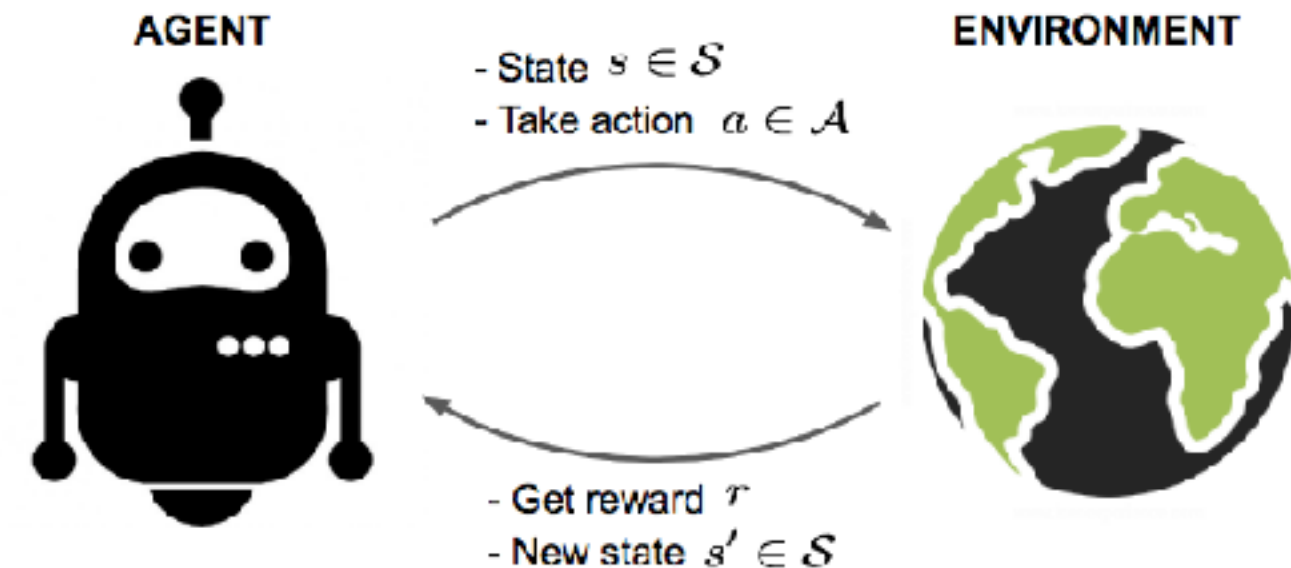
Unsupervised Learning



- Learning patterns / structures from **unlabeled data**
- Goal : identify interesting **patterns** in the data
- Applications : clustering, association
- Example algorithms : PCA
AutoEncoder
- Challenges : evaluating whether the algorithm is learning something useful

Types of Machine Learning

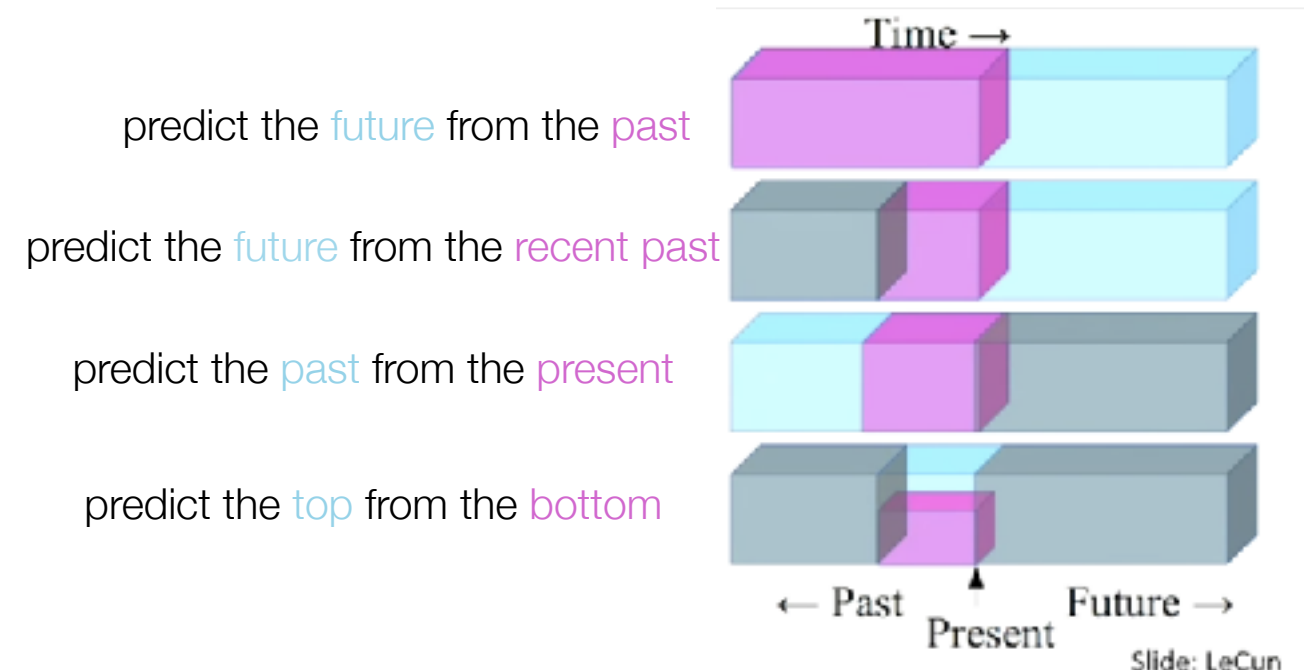
Reinforcement Learning



- Learning by trial and error
- Goal : maximizing the reward
- Example algorithms : AlphaZero
(trained entirely from self-play)

Self-Supervised Learning

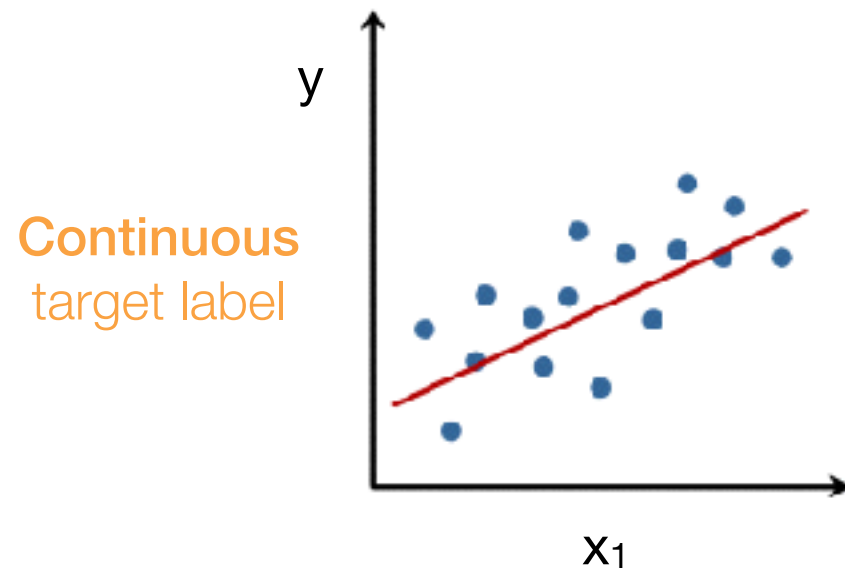
Machine learns to predict part of the input from any other part.



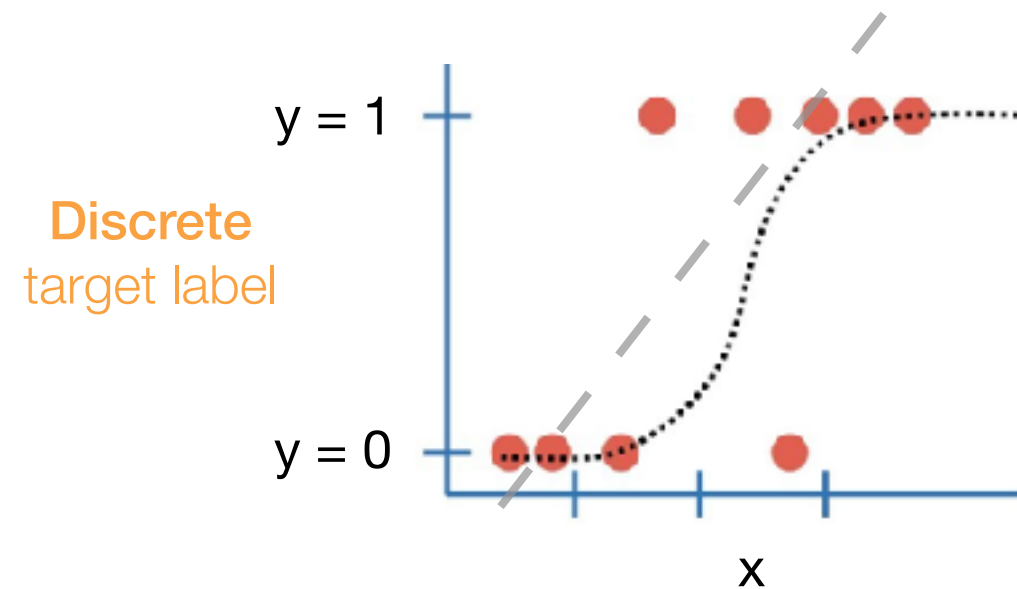
- Learning patterns by better utilizing **unlabeled data**.
- Goal : Learn **intermediate representations with good structural meanings** and can be beneficial to a variety of practical downstream tasks.
- “Most of what we learn as humans is in a self-supervised mode, not a reinforcement mode. It’s basically **observing the world and interacting with it a little bit**, mostly by observation in a test-independent way.” — Yann LeCun

Classical Machine Learning Methods

Linear Regression



Logistic Regression



model prediction

$$\hat{y} \approx w_0 + w_1 x_1$$

approximately modeled as

logistic function

$$\hat{y} = p(y=1 | x) \approx \frac{1}{1 + e^{-(w_0 + w_1 x)}}$$

↖ logit / log-odd

$$\begin{aligned} \log\left(\frac{\text{prob. of classifying to 1}}{\text{prob. of classifying to 0}}\right) &= \log\left(\frac{p(y=1|x)}{p(y=0|x)}\right) \\ &= \log\left(\frac{\hat{y}}{1-\hat{y}}\right) = w_0 + w_1 x \end{aligned}$$

Classical Machine Learning Methods

Linear Regression

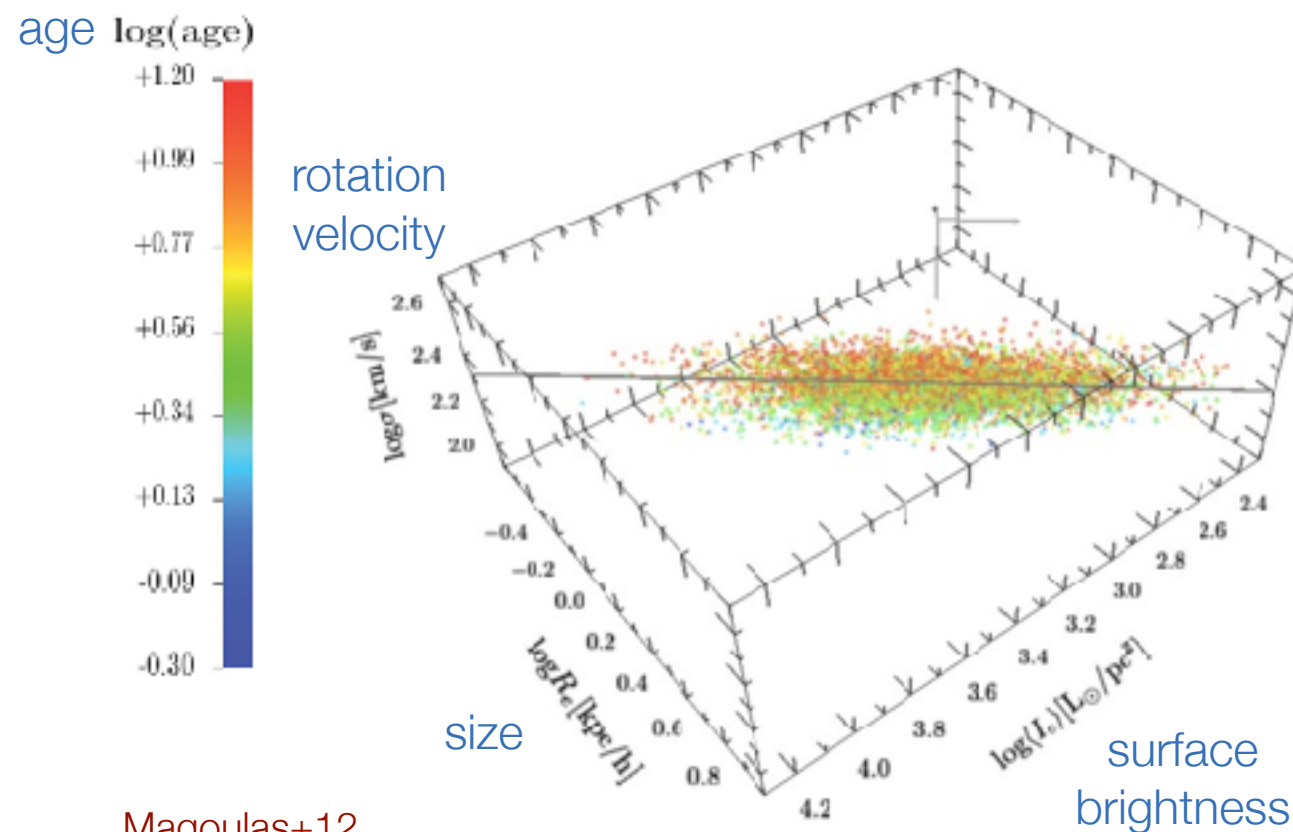
$$\hat{y} \approx w_0 + \sum_i w_i x_i$$

Logistic Regression

$$\hat{y} \approx \frac{1}{1 + e^{-(w_0 + \sum_i w_i x_i)}}$$

High Dimensional feature space

Fundamental Plane of galaxies

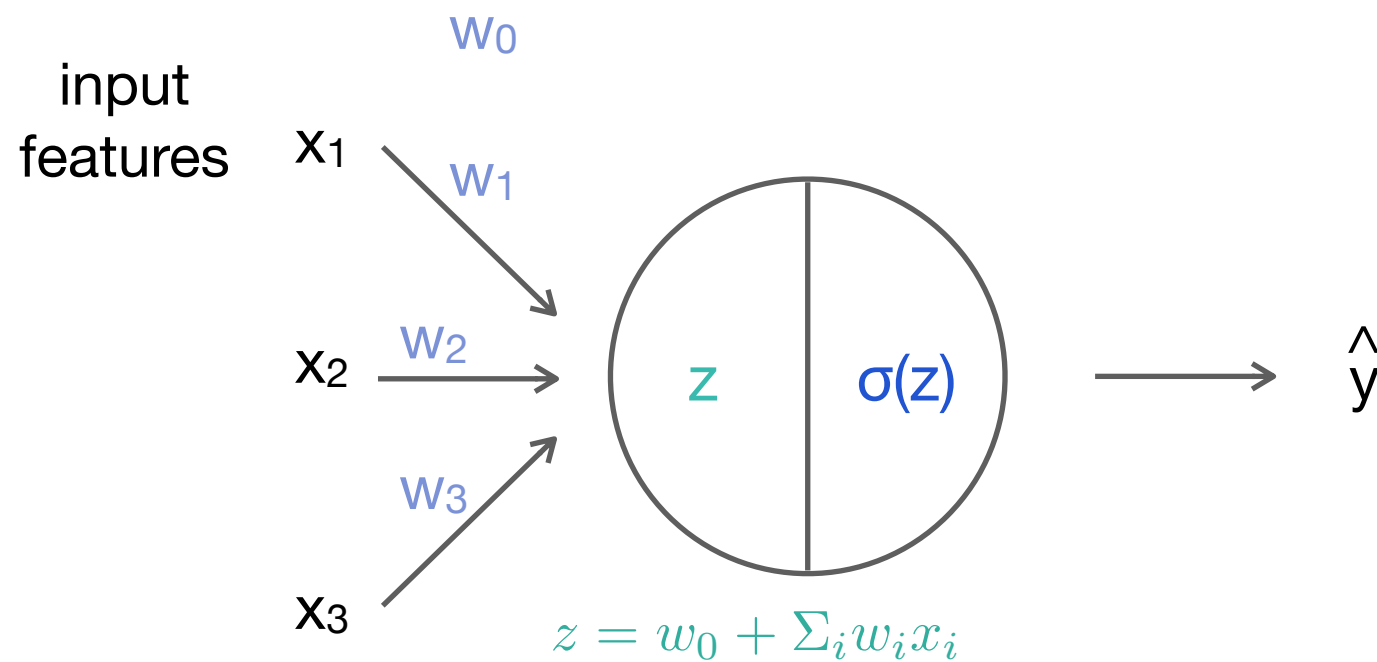


Magoulas+12

Active galaxy ? $y = 1$: active galaxy
 $y = 0$: quiescent galaxy

stellar mass
emission line intensity
black hole mass
color
...

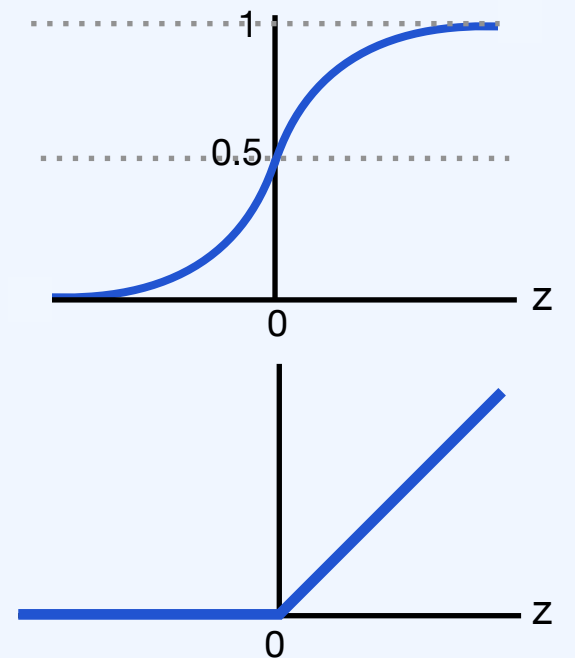
Neural Network – Single Neuron



activation functions

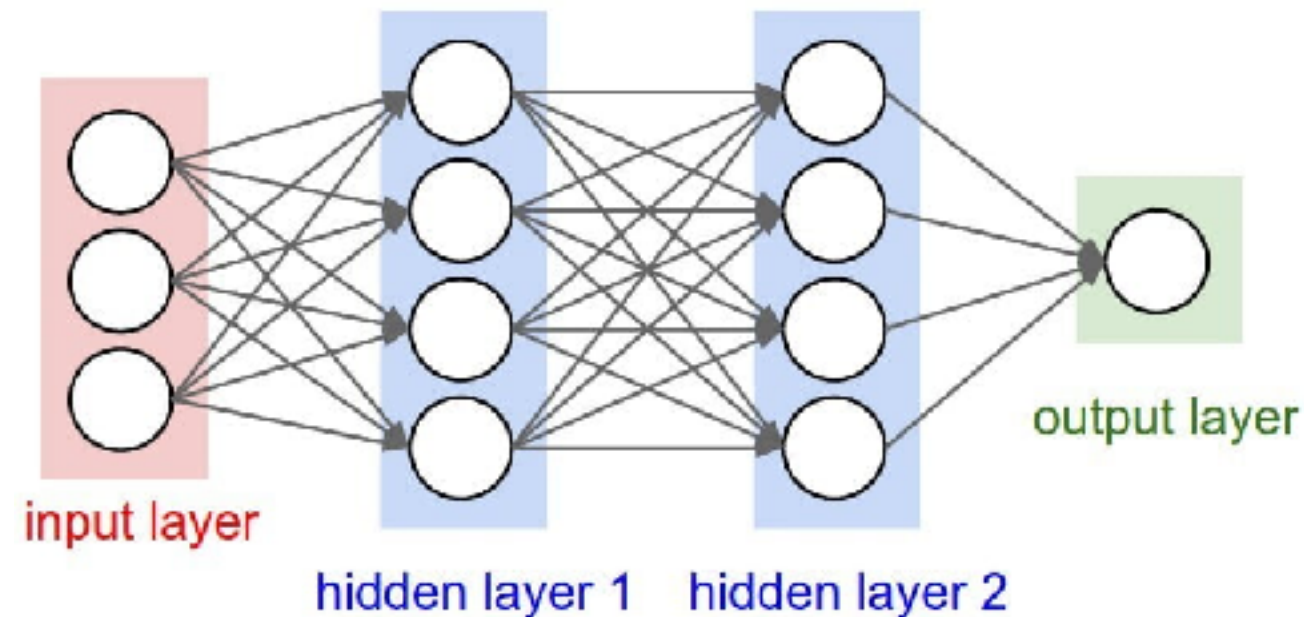
sigmoid : $\sigma(z) = \frac{1}{1 + e^{-z}}$

ReLU : $\sigma(z) = \max(0, z)$



- Activation functions add non-linearity to neural network models.

Neural Network – Fully connected Multilayer Perceptrons (MLP)

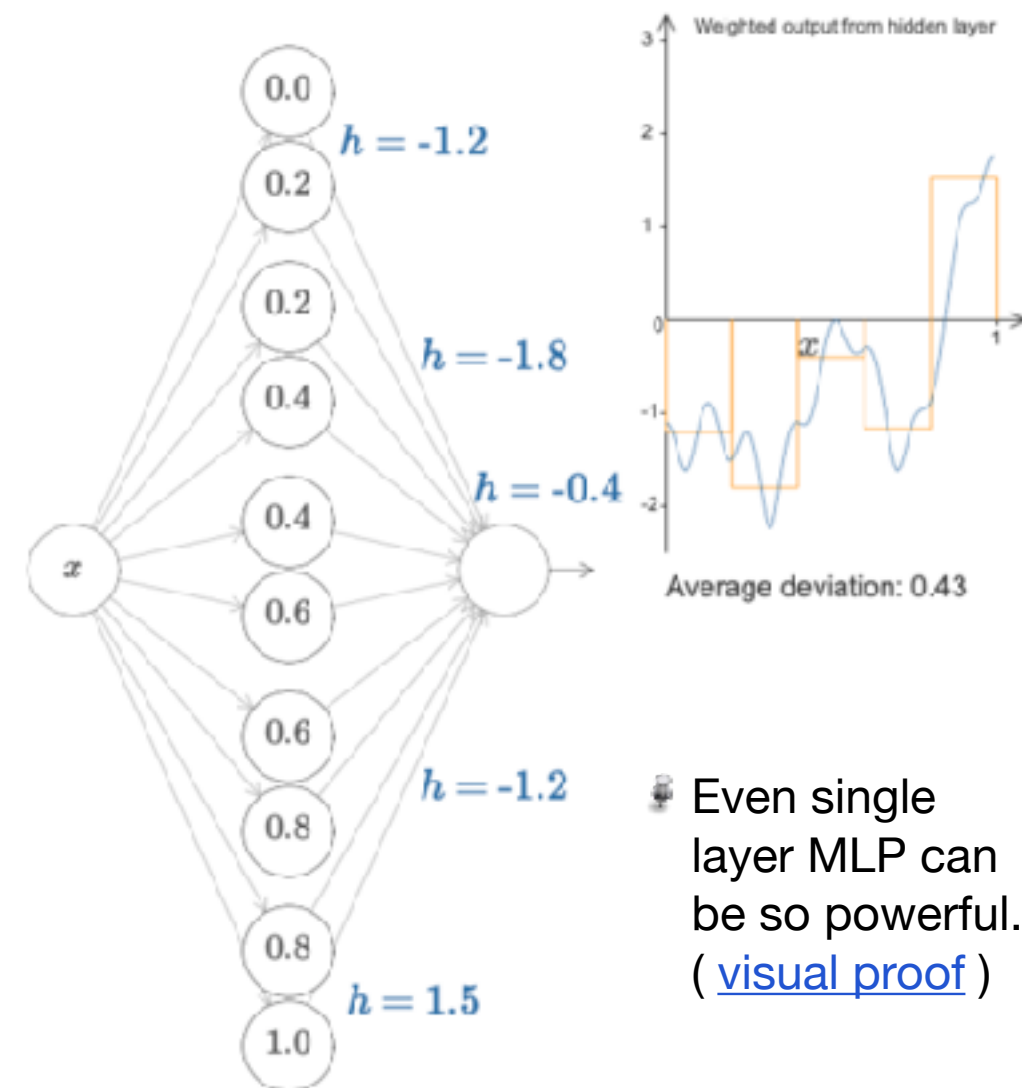


Universal Function Approximation Theorem

- By growing the network size, MLPs can approximate **any continuous functions** up to the **desired accuracy level**.

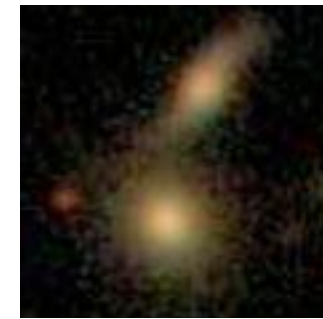
Width v.s Depth

- Deep networks can learn features in a **hierarchical** way
Earlier layers : simple structures like edges
Deeper layers : more complex representations
- Wide, shallow networks are more likely to have **overfitting** issue.



- Even single layer MLP can be so powerful. ([visual proof](#))

Neural Network – Fully connected Multilayer Perceptrons (MLP)



128 x 128 pixels

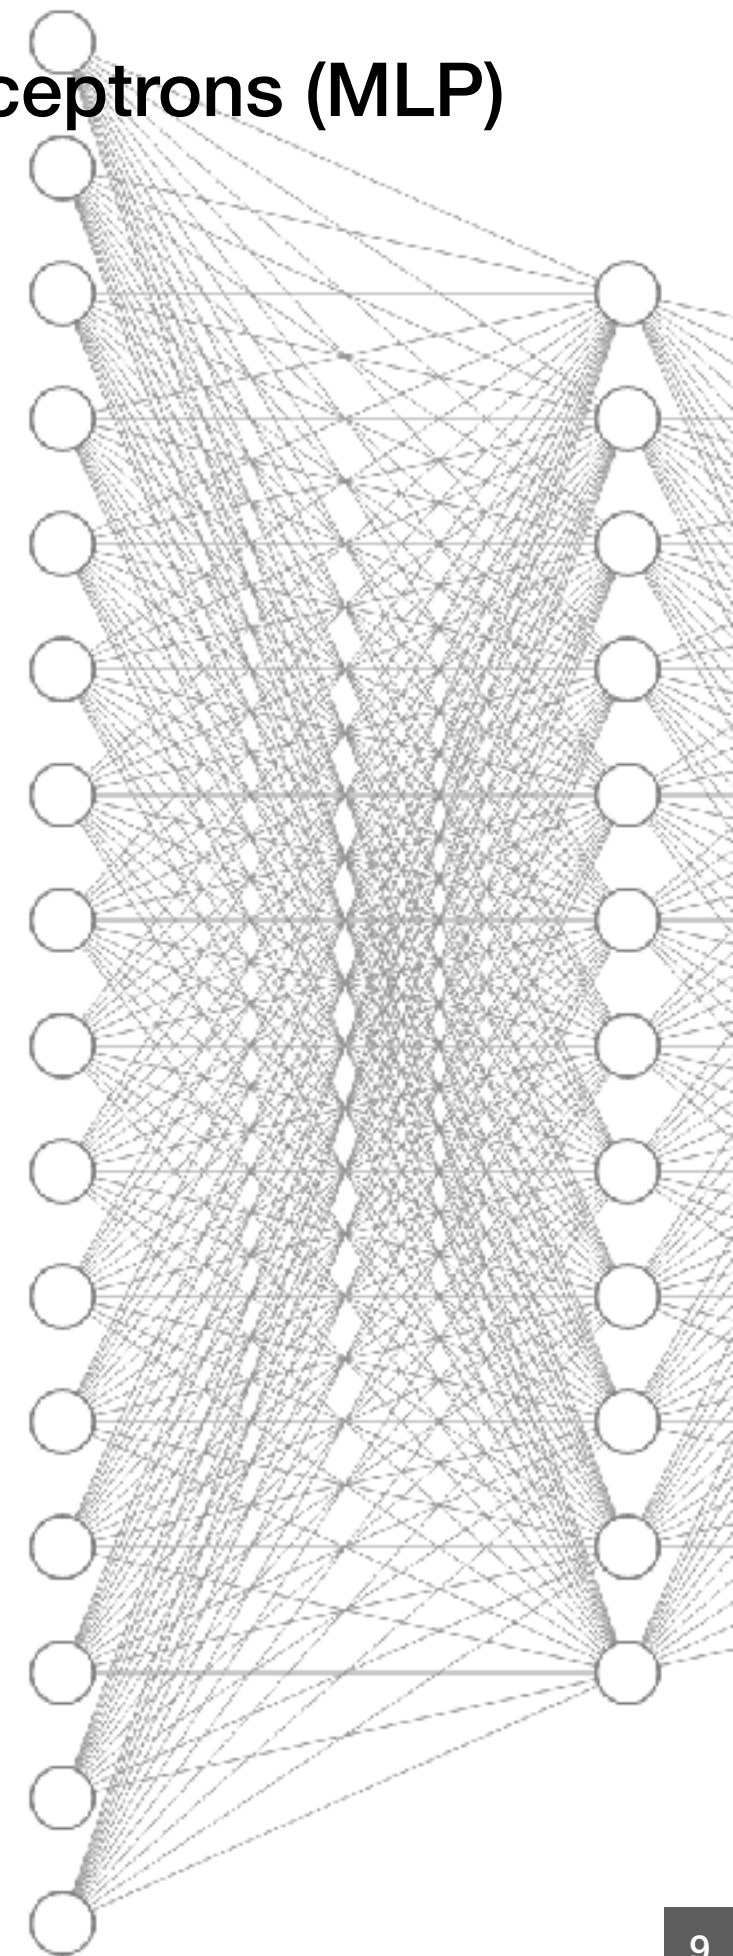
Limitations of MLP

- **Too many parameters** due to fully connection.
- Low data efficiency : need lots of data to learn well.

Fully connected structure



No prior assumption on how features interact from data.



Neural Network – Convolutional Neural Network (CNN)

Image Convolution

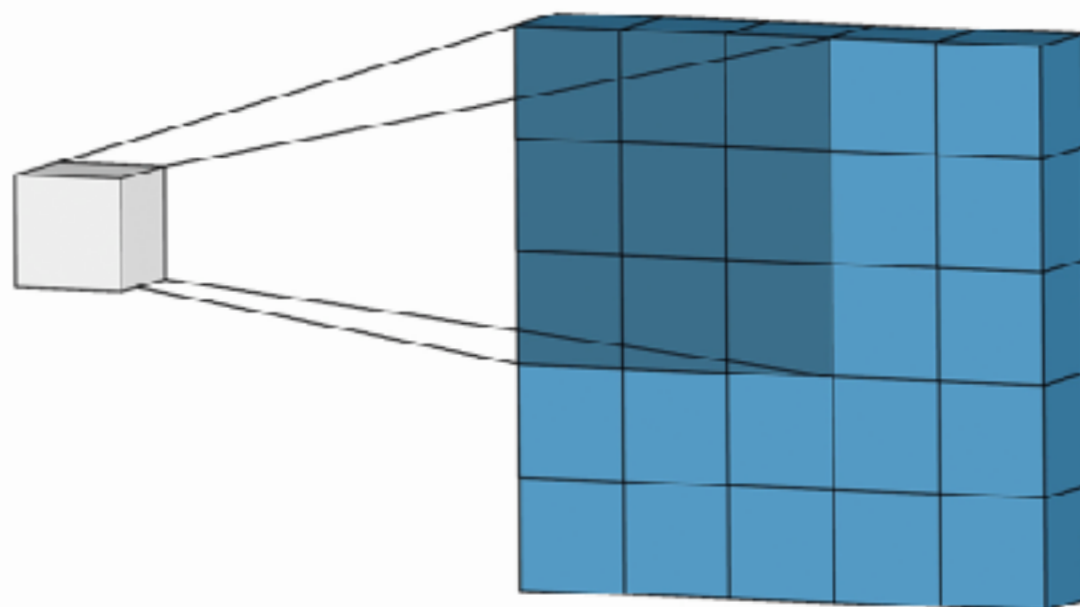
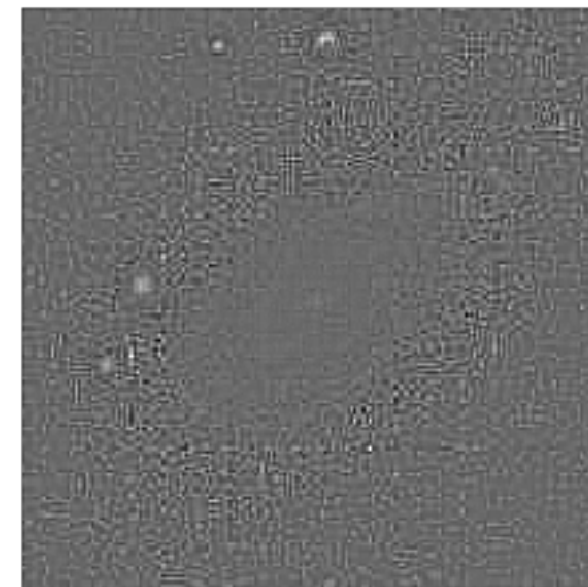


filter / kernel

$$\begin{matrix} * & \begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix} & = \end{matrix}$$

learnable parameters

feature map



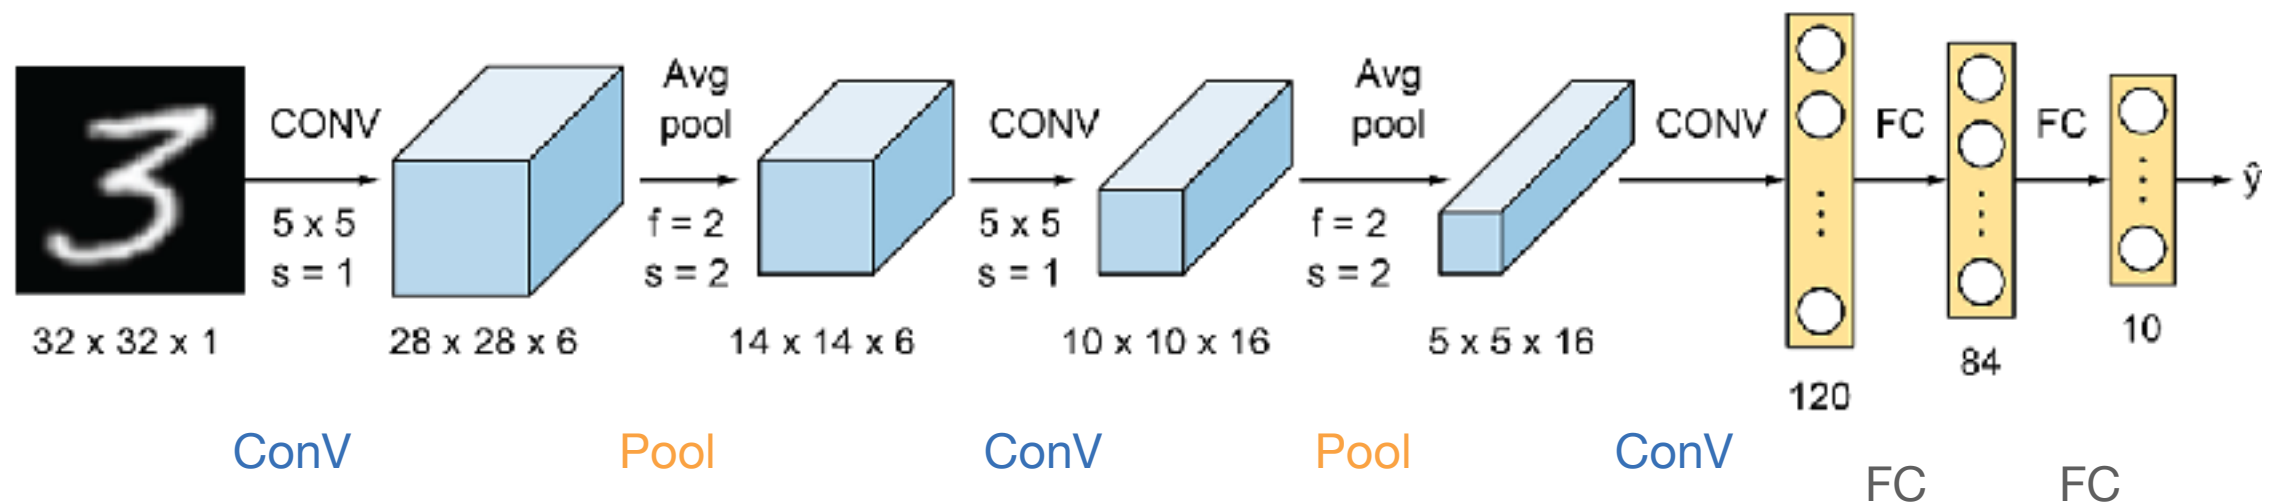
CNN Properties

- **Translational Symmetry**
Weight sharing across the entire image
- **Local Connectivity**
Drop connections between far away neurons

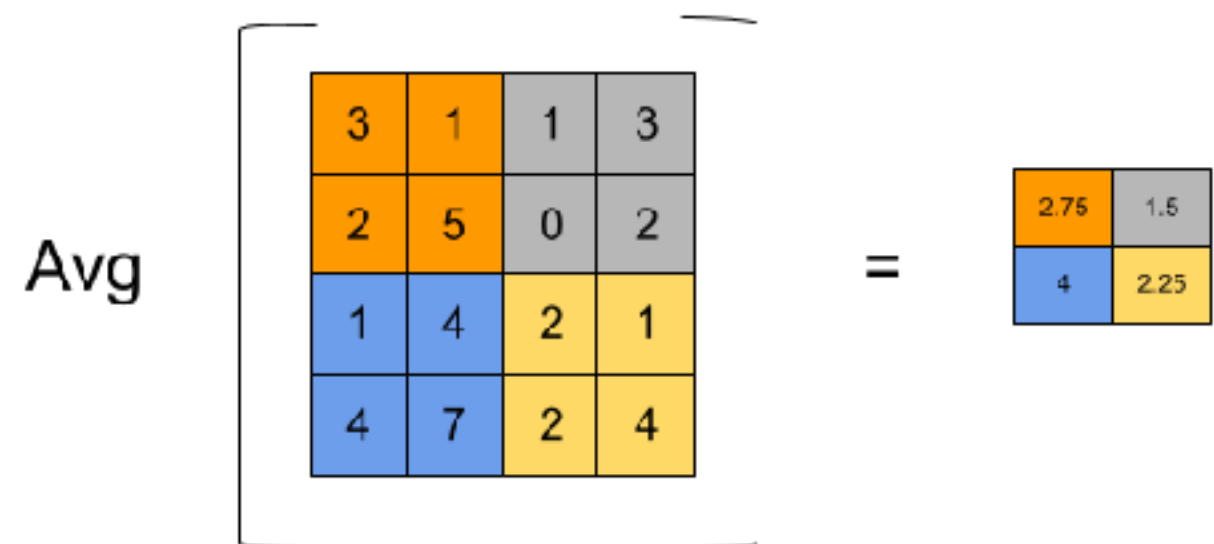
Convolutional Network structures

LeNet-5 ~ 60k parameters

A pioneering CNN network by LeCun et al. 1998



Average Pooling Layer

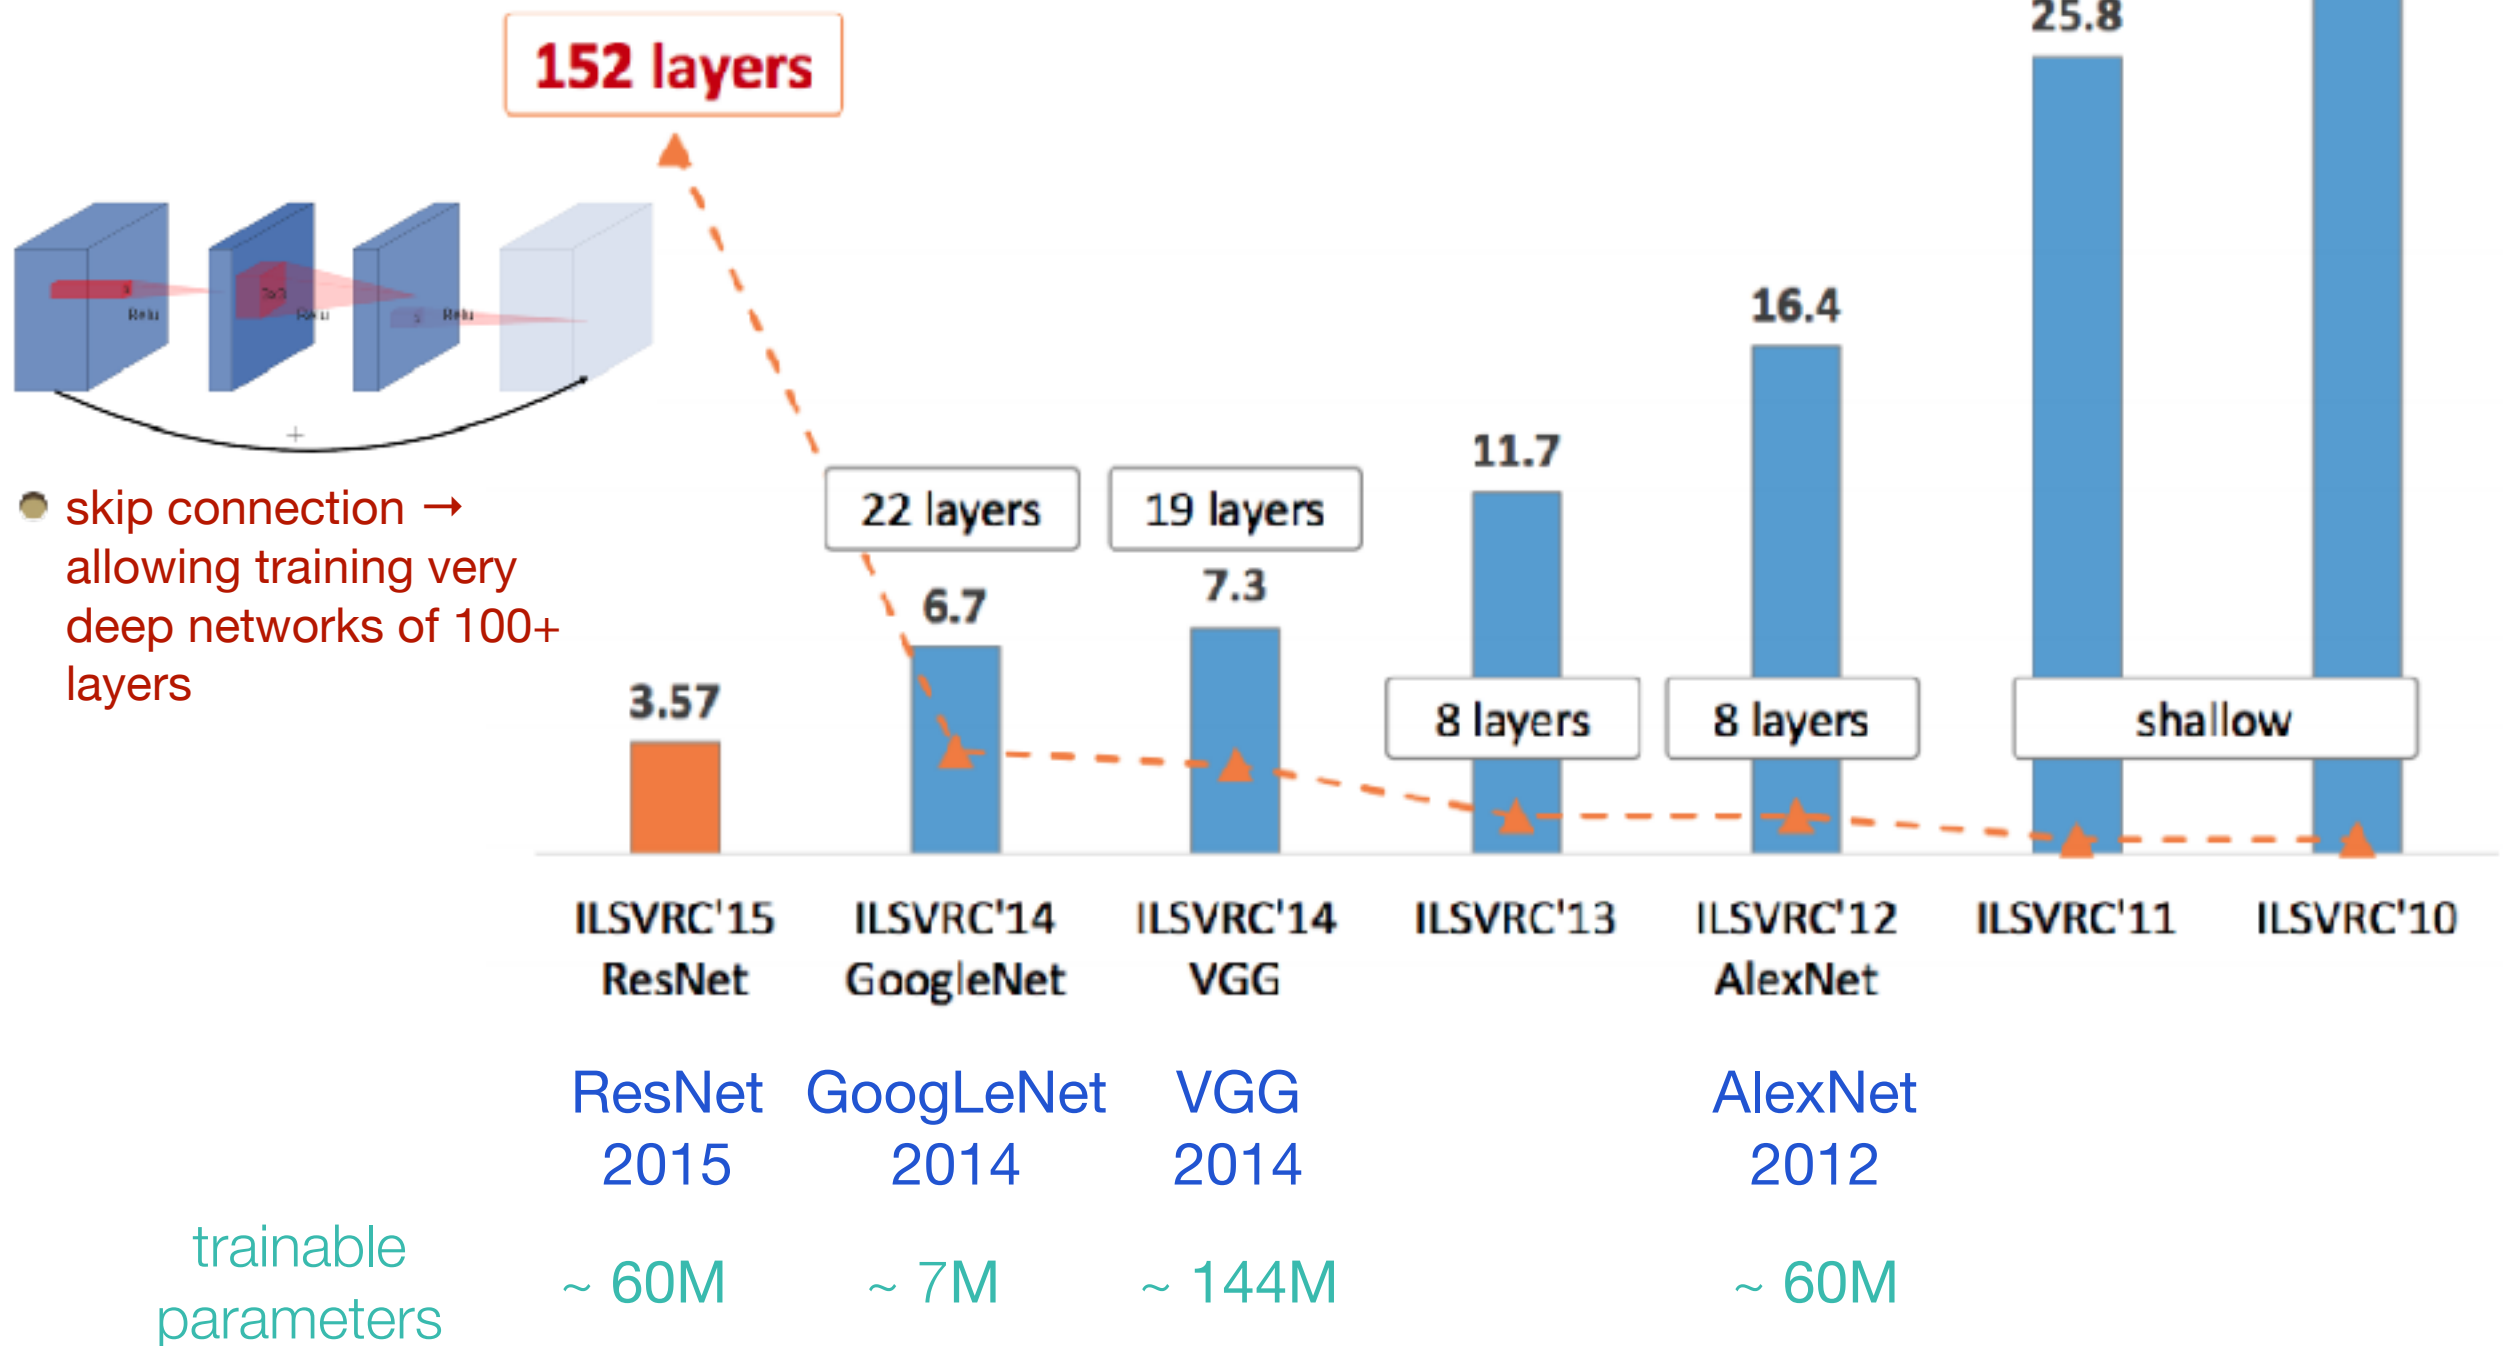


Evolution of CNNs

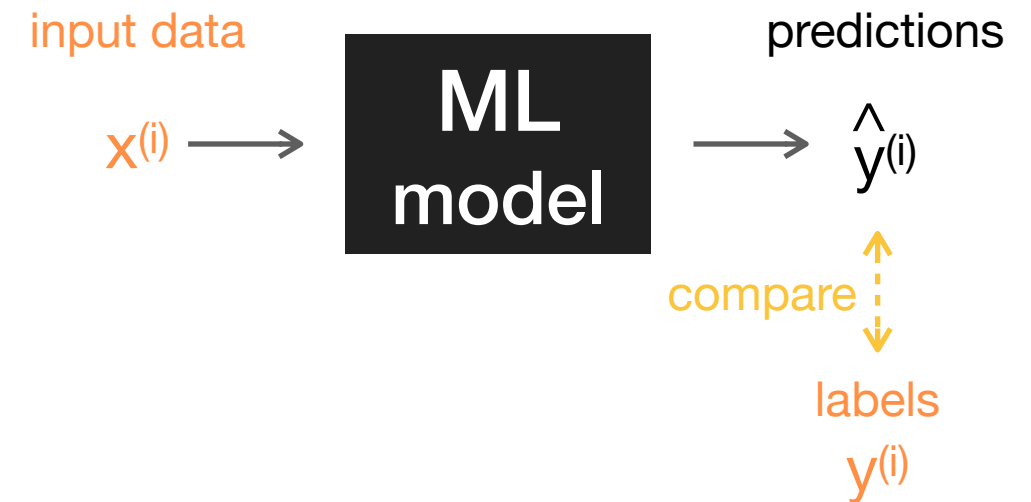
1000 categories

if the target label is one of the top 5 predictions
→ correct prediction

Classification: ImageNet Challenge top-5 error



Training Process



0. Training Data

$(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots (x^{(i)}, y^{(i)}) \dots, (x^{(m)}, y^{(m)})$

1. Define Model

- From simple \rightarrow more complex network structures.
- For similar data types \rightarrow Find existing network model and apply directly.

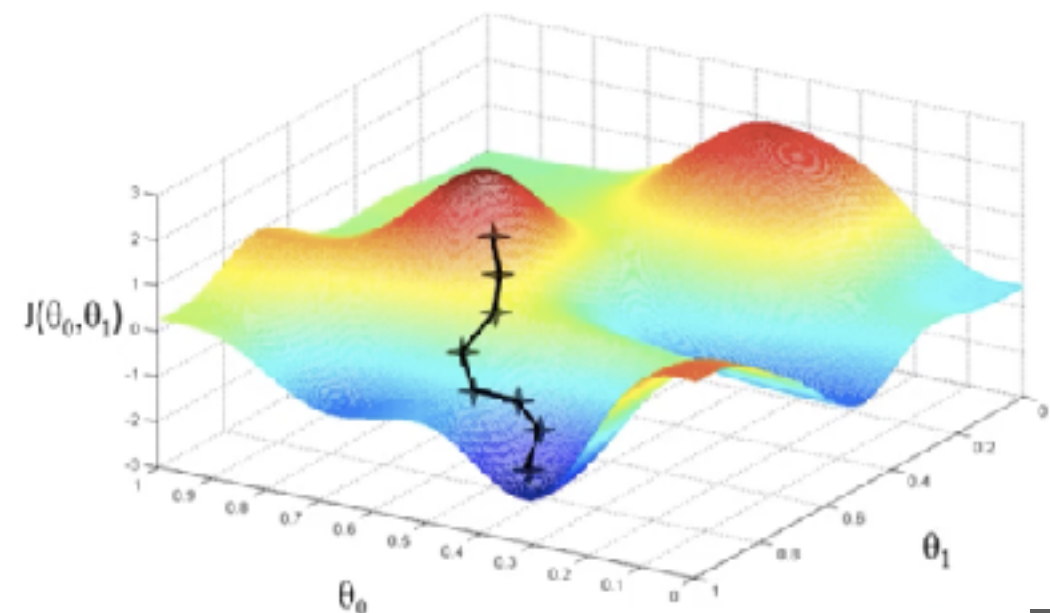
2. Define Loss

- Mean Squared Loss : $\frac{1}{2}(\hat{y} - y)^2$
 - Cross Entropy Loss : $-(y \log \hat{y} + (1 - y) \log(1 - \hat{y}))$
- average across all training samples \longrightarrow
- $$J = \frac{1}{2m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2$$
- $$J = \frac{-1}{m} \sum_{i=1}^m (y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}))$$

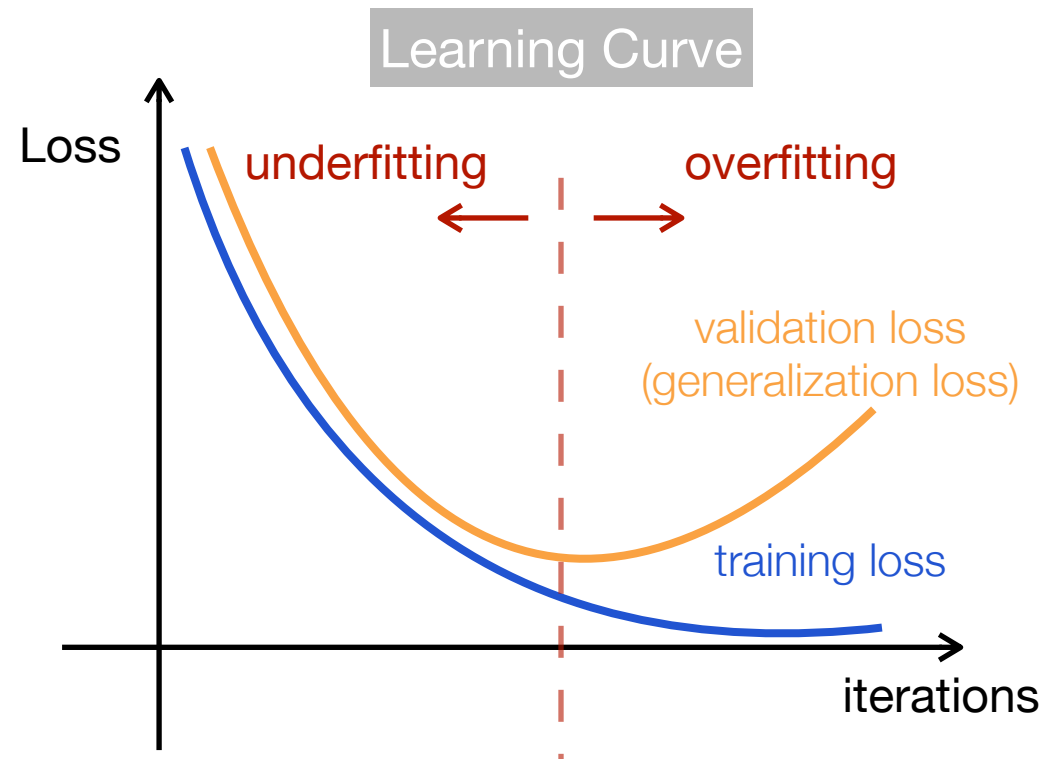
3. Optimization

α : learning rate

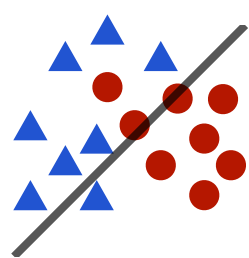
- Gradient Descent $\theta_j = \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$
- Advanced optimization techniques / algorithms :
Momentum, RMSProp, Adam, Learning rate decay



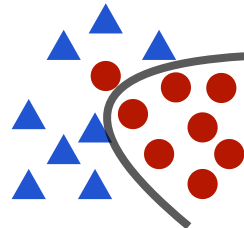
Overfitting & Regularization Techniques



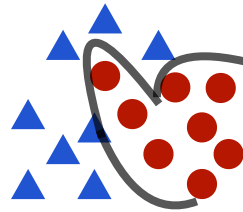
● Bias & Variance



high bias



just right



high variance

- **Regularization** : techniques that discouraging learning a more complex model to prevent overfitting.

Data augmentation

natural image : random shift, color change
galaxy image : random rotation, foreground contamination

Early stopping

Select the model that **performs the best** on the **validation set**.

L2 regularization

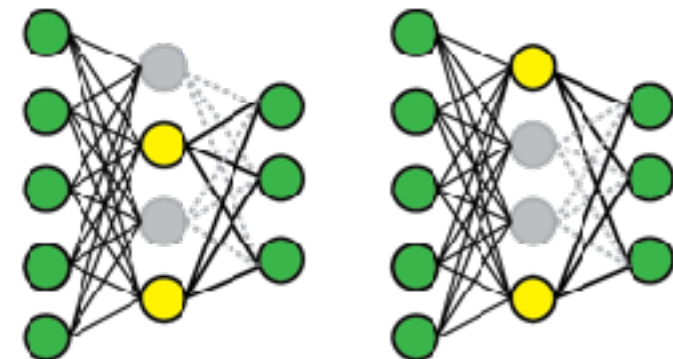
minimize (**Loss(Data|Model)** + **Complexity(Model)**)

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y) + \frac{\lambda}{2m} \|\theta\|^2$$

λ : regularization parameter

Dropout

Randomly eliminate a fraction of nodes for each training example.



Summary

Classical Machine Learning

- Works better for **structured data** (catalog, tabular data).
- Limited performance on learning complex functions.
- Data Efficient. Easier to Train.
- Require some data manipulation/exploration before feeding to the algorithm (dimensionality reduction, feature extraction...).
- Easier to interpret.
- Usually with evaluable prediction uncertainty.
- Not covered in this course. But is really useful in physical science.
 - ▶ Statistics, Data Mining, and Machine Learning in Astronomy — Ivezić et al.
 - ▶ [ASTR 502 2020 class notebook](#)

Deep Learning

- Works especially well for **unstructured data** (e.g. image, audio signal).
- Superior performance on wide variety of tasks.
- Data efficiency is poor. Need lots of data to train.
- Directly pass the data into the network.
- Difficult to understand.
- Challenge to evaluate uncertainty.
- Focus of this course.
 - ▶ [ML papers in cosmology](#)

