

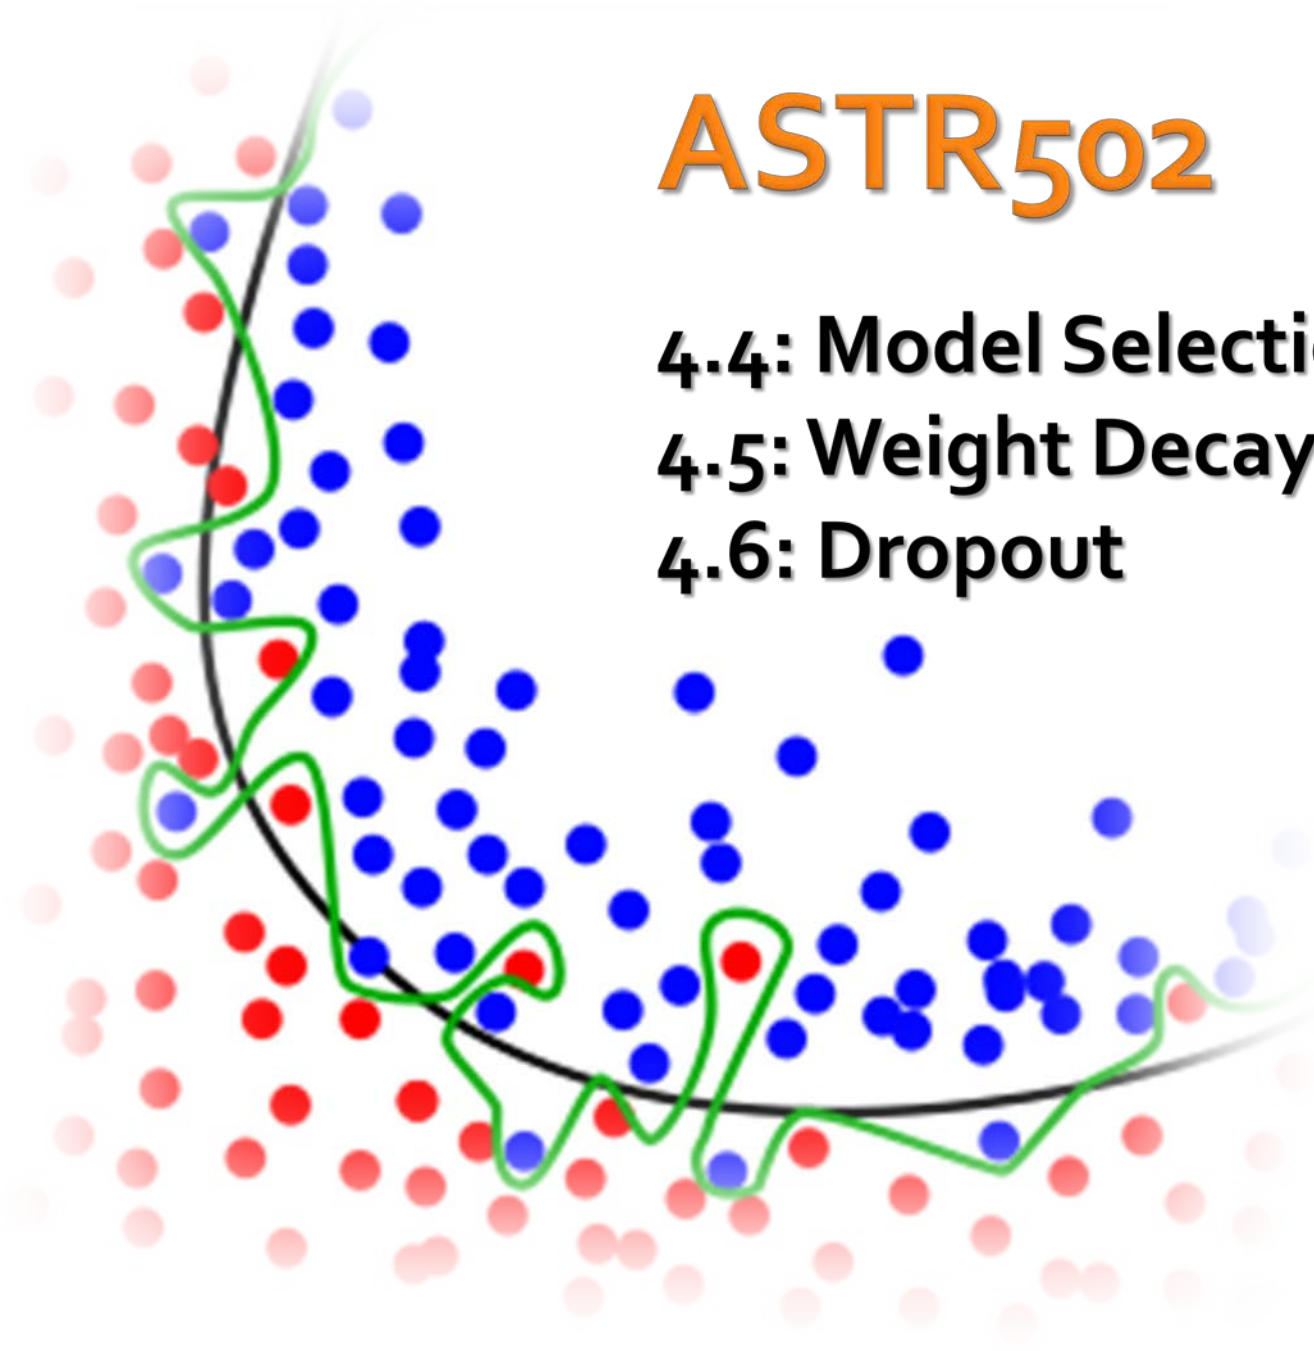
ASTR502

4.4: Model Selection, Underfitting, Overfitting

4.5: Weight Decay

4.6: Dropout

Donghyeon Jeff Khim



The fundamental problem (and goal) of ML

- Discovering general patterns
 - General pattern vs. (Simply) memorized data
 - Our predictions will only be useful if our model has truly discovered a **general pattern**

Training data

DB
ML

Test (Trained)

Test (New)

Types of Galaxies



DB - Result: Spiral
ML - Result: Spiral



DB - Result: Not found
ML - Result: Spiral

The fundamental problem (and goal) of ML


- Discovering general patterns
 - General pattern vs. (Simply) memorized data
 - Our predictions will only be useful if our model has truly discovered a **general pattern**

How much we can trust the results
from the machine learning?

Training data

Test (New)

Types of Galaxies



DB - Result: Spiral
ML - Result: Spiral




DB - Result: Not found
ML - Result: Spiral

Training Error & Generalization Error

- **Training Error:** the error calculated on the training dataset
- **Generalization Error:** model's error for an infinite amount of data
 - We can never calculate the generalization error exactly → **Expectation**

Training data

Types of Galaxies



Training Error
Accuracy : **99%**
(Measured by
training sets)

Real data



Generalization Error
Accuracy : **99%**
(Estimated)


ML - Result: Spiral

Training Error & Generalization Error

- **Training Error:** the error calculated on the training dataset
- **Generalization Error:** model's error for an infinite amount of data
 - We can never calculate the generalization error exactly → **Expectation**

Training data

Types of Galaxies



Training Error
Accuracy : **99%**
(Measured by
training sets)

Overfitting!




Real data



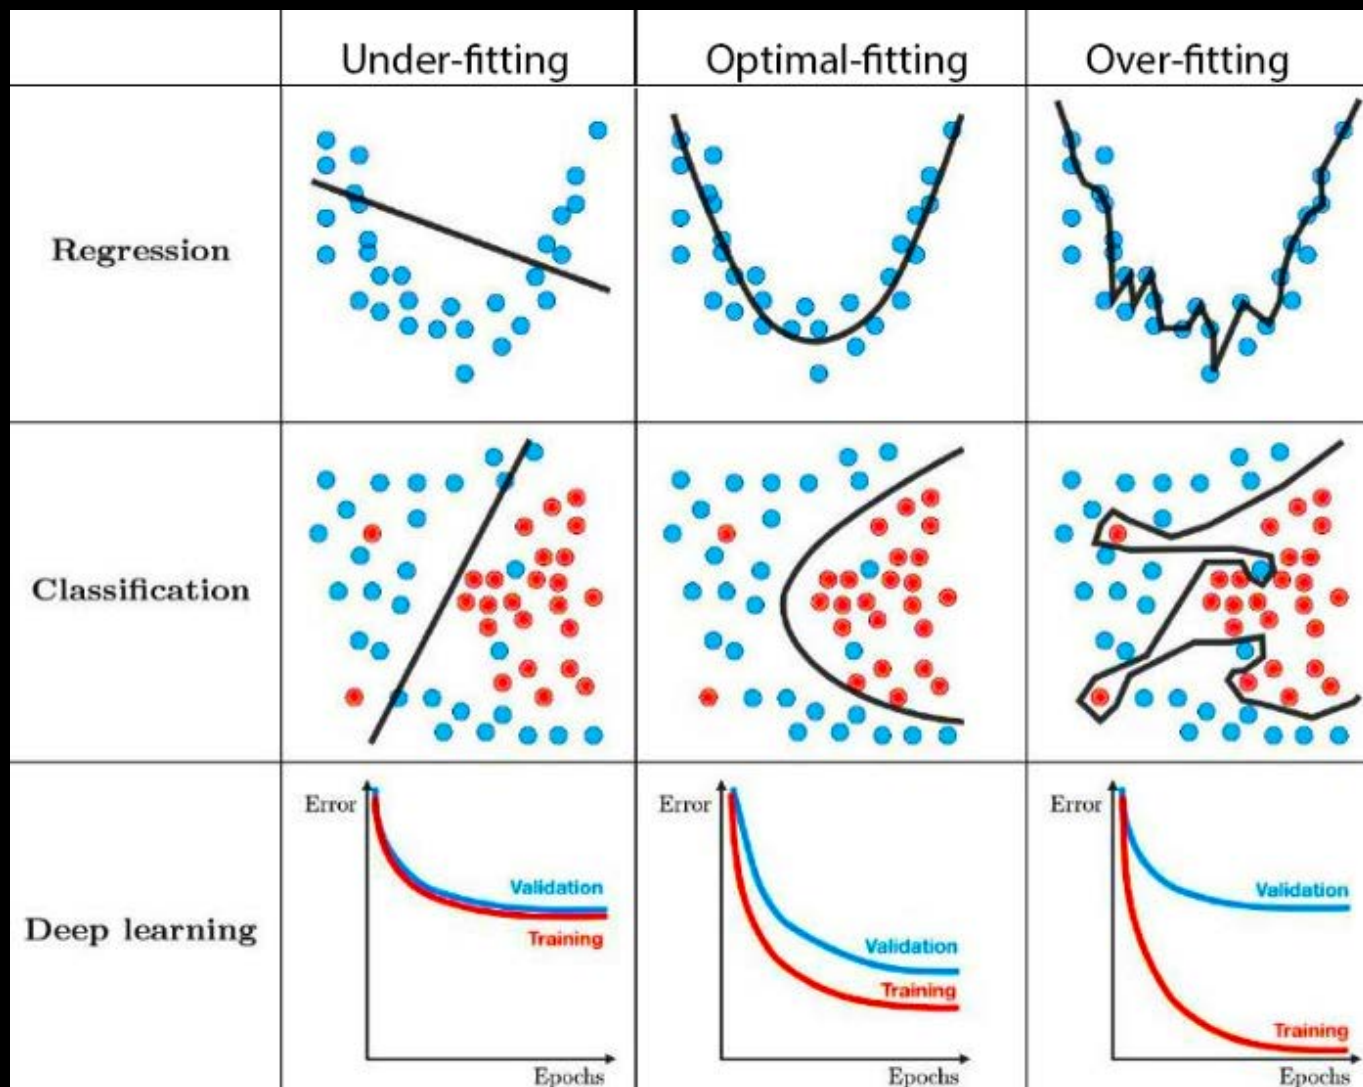
Generalization Error
Accuracy : **80%**
(Estimated)

ML - Result: Spiral

Underfitting and Overfitting (Regression)



Underfitting and Overfitting



Generalization Error
Accuracy : **80%**
(Estimated)

Real data



Training data

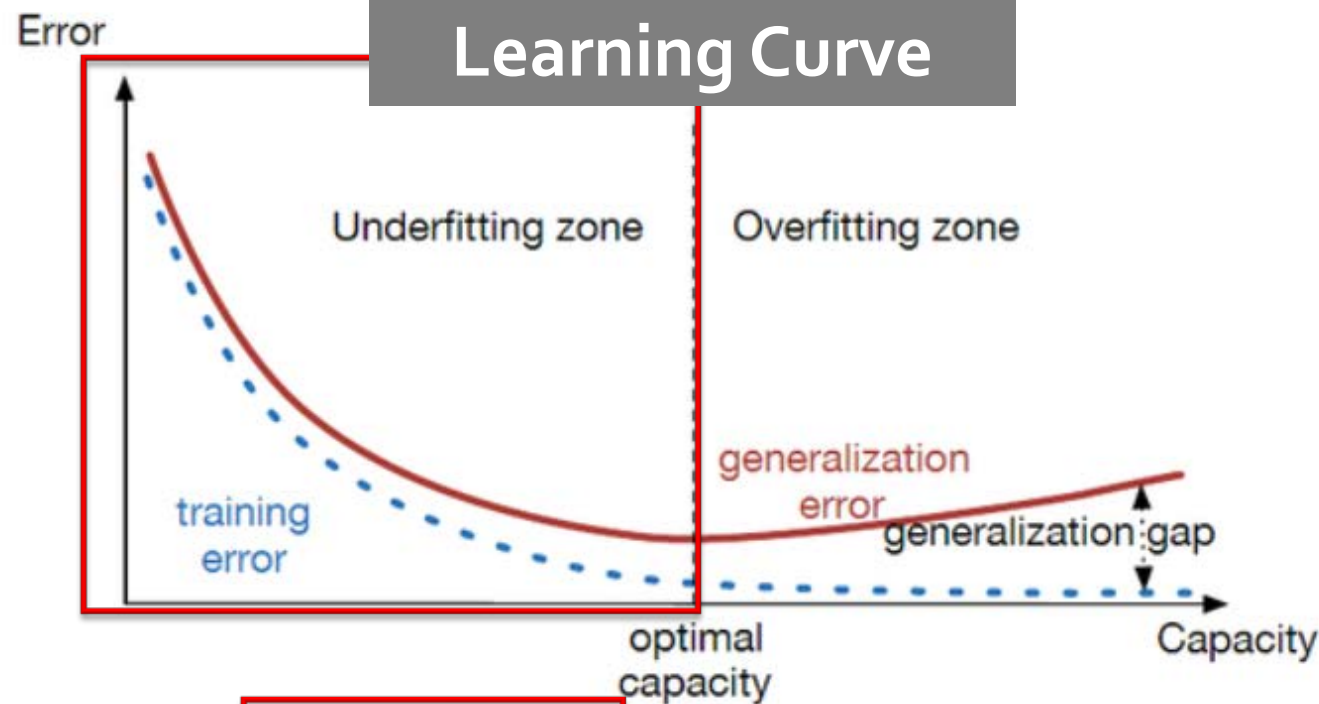
Types of Galaxies




Training Error
Accuracy : **99%**
(Measured by
training sets)

Underfitting and Overfitting

Learning Curve



Deep learning

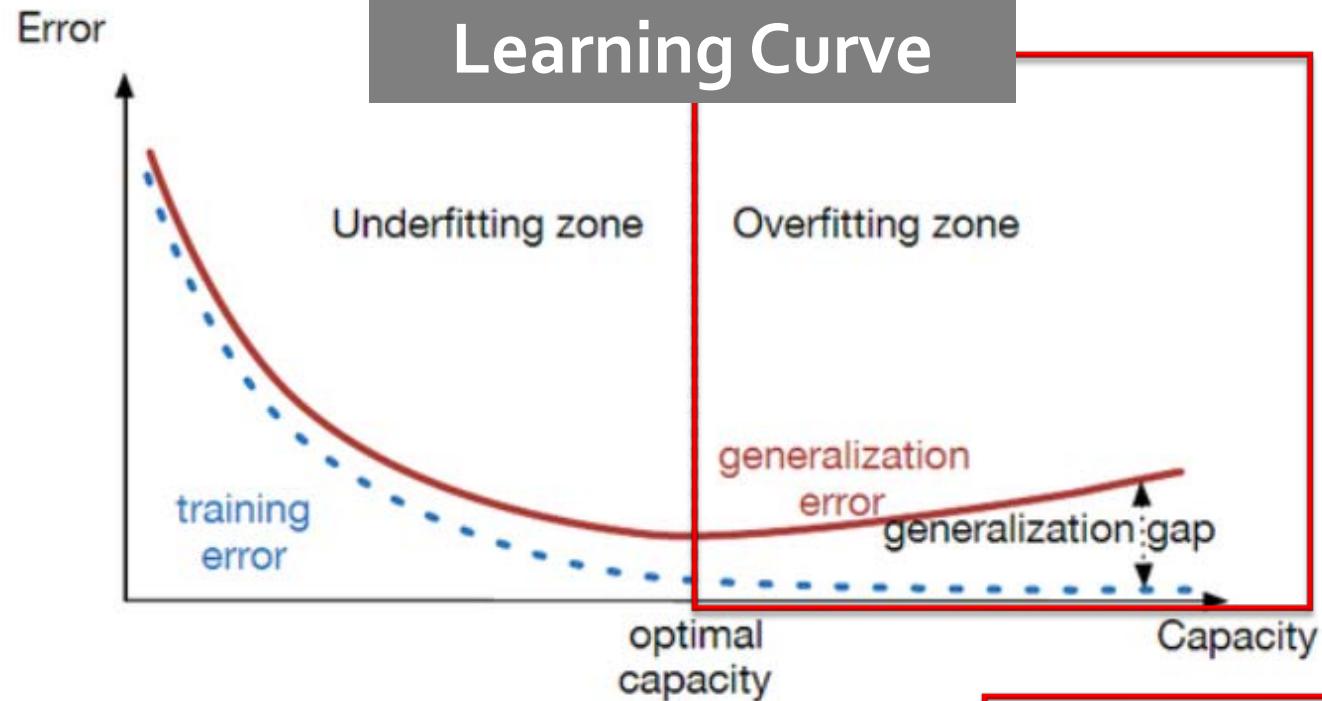


■ Underfitting

- Training & validation errors are both substantial
 - The model is too simple to capture the pattern
- Small generalization gap
 - We may use more complex models
- **High bias**
 - Not be able to fit data well

Underfitting and Overfitting

Learning Curve




■ Overfitting

- Training error is significantly lower than the validation error
 - Huge generalization gap

■ High variance

- Performs specifically well under certain noise realization of data

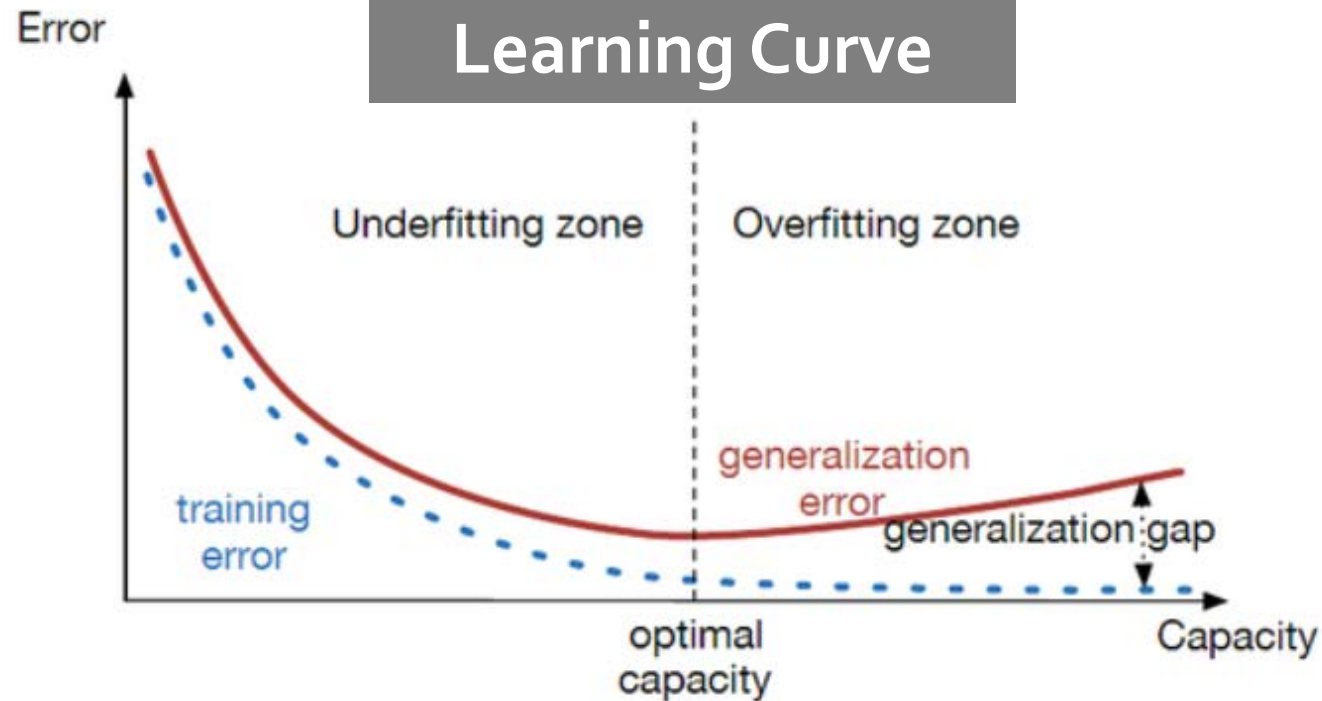
Deep learning



Underfitting and Overfitting

Page 4

Learning Curve



- **Training Error**: the error calculated on the training dataset
- **Generalization Error**: model's error for an infinite amount of data
 - We can never calculate the generalization error exactly → **Expectation**

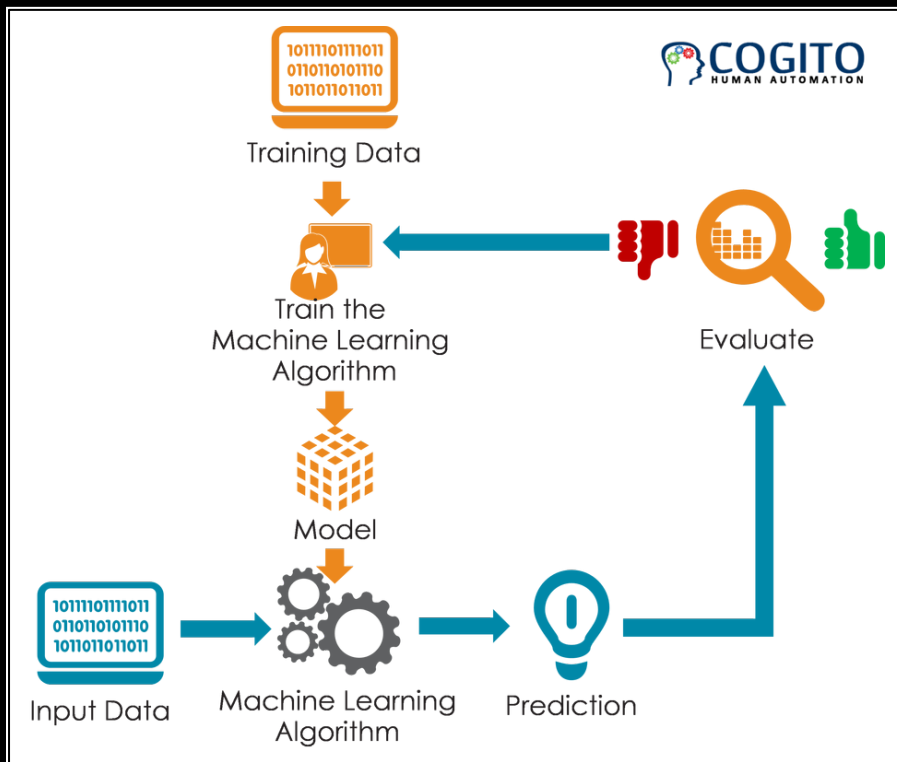
How can we measure (estimate) the errors?

Training set, Test set

Data set (with labels)

Training set

Test set



■ Training Set

- Dataset that we use to train the model to determine the network parameters (weights and biases)

■ Test Set

- Evaluate the network
- Provide an **unbiased estimate** on the performance of the final network.
- We *may* measure the validation error based on the test set


Can we say the test set did not affect the model?

Training set, Test set

Data set (with labels)

Training set

Test set



Check Overfitting, Underfitting
Hyperparameters - # of layers, # of
hidden units per layer, Batch size,
learning rate...

■ Test set

- We should not touch our test set until after we have chosen all our hyperparameters.
- There is a risk that we **overfit** the test data
- → We should never rely on the test data for model selection.

Training set, Test set & Validation set

Data set (with labels)

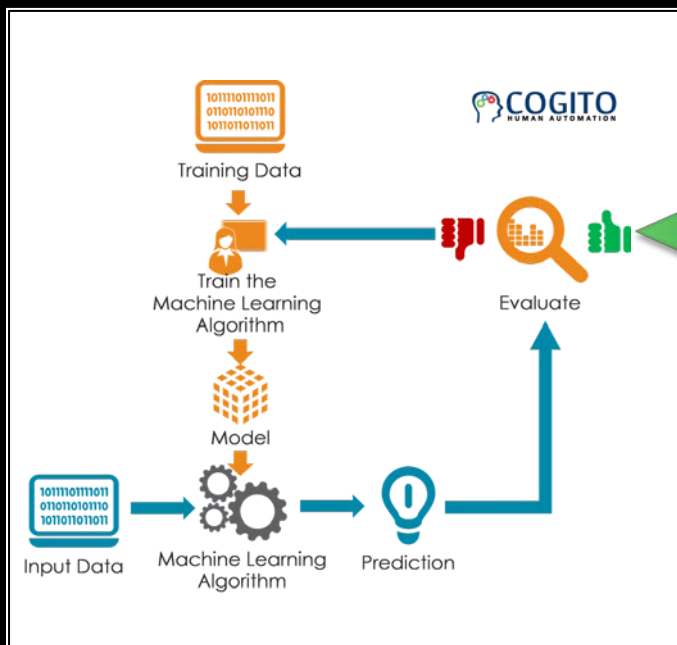
Training set

Test set

Training set

Validation set

Test set



Validate model

Final evaluation

→ Unbiased estimate on the performance

Test set

- We should not touch our test set until after we have chosen all our hyperparameters.
- There is a risk that we **overfit** the test data
- → We should never rely on the test data for model selection.

Validation set

- Validate model performance during training.
- Provide information on the generalizability of our model to unseen data
- Helpful to prevent overfitting

K-Fold Cross-Validation

Data set (with labels)

Training set

Test set

Training set

Validation set

Test set

Fold 1

Fold 2

Fold 3

Fold 4

Fold 5

Fold 1

Fold 2

Fold 3

Fold 4

Fold 5

Fold 1

Fold 2

Fold 3

Fold 4

Fold 5

Fold 1

Fold 2

Fold 3

Fold 4

Fold 5

Fold 1

Fold 2

Fold 3

Fold 4

Fold 5

Finding
parameters

Test set

Final evaluation

$$Error = \frac{1}{5} \sum Err_i$$

- When training data is **scarce**
 - We might not even be able to afford to hold out enough data to constitute a proper validation set.
- K-fold cross-validation
 - The original training data is split into K non-overlapping subsets.
 - Model training and validation are executed K times
 - Training and validation errors are estimated by averaging K errors

Model Complexity

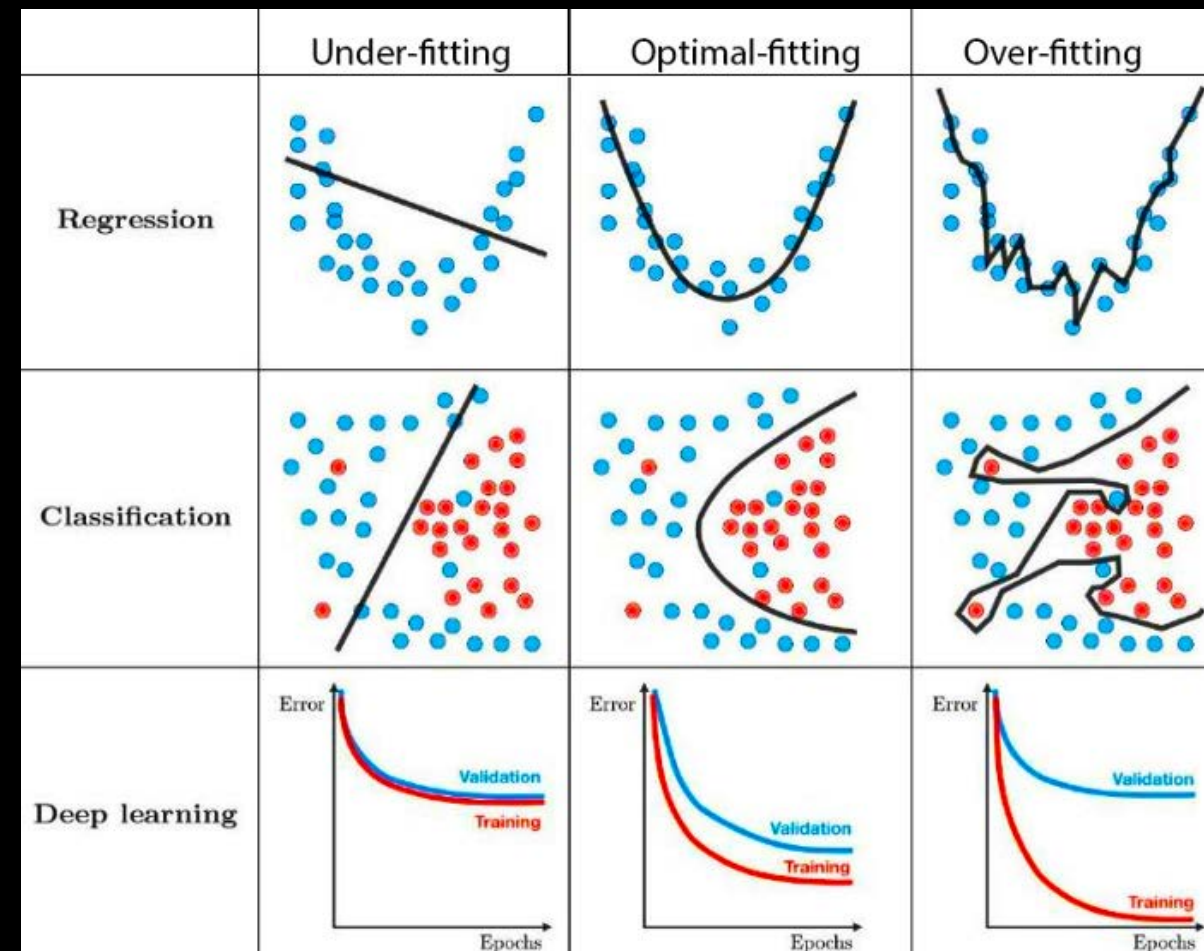
■ Complexity & Fitting

■ Underfitting

- Simple models and abundant data
- Generalization error \cong training error

■ Overfitting

- More complex models and smaller data
- Training error \downarrow , generalization gap \uparrow

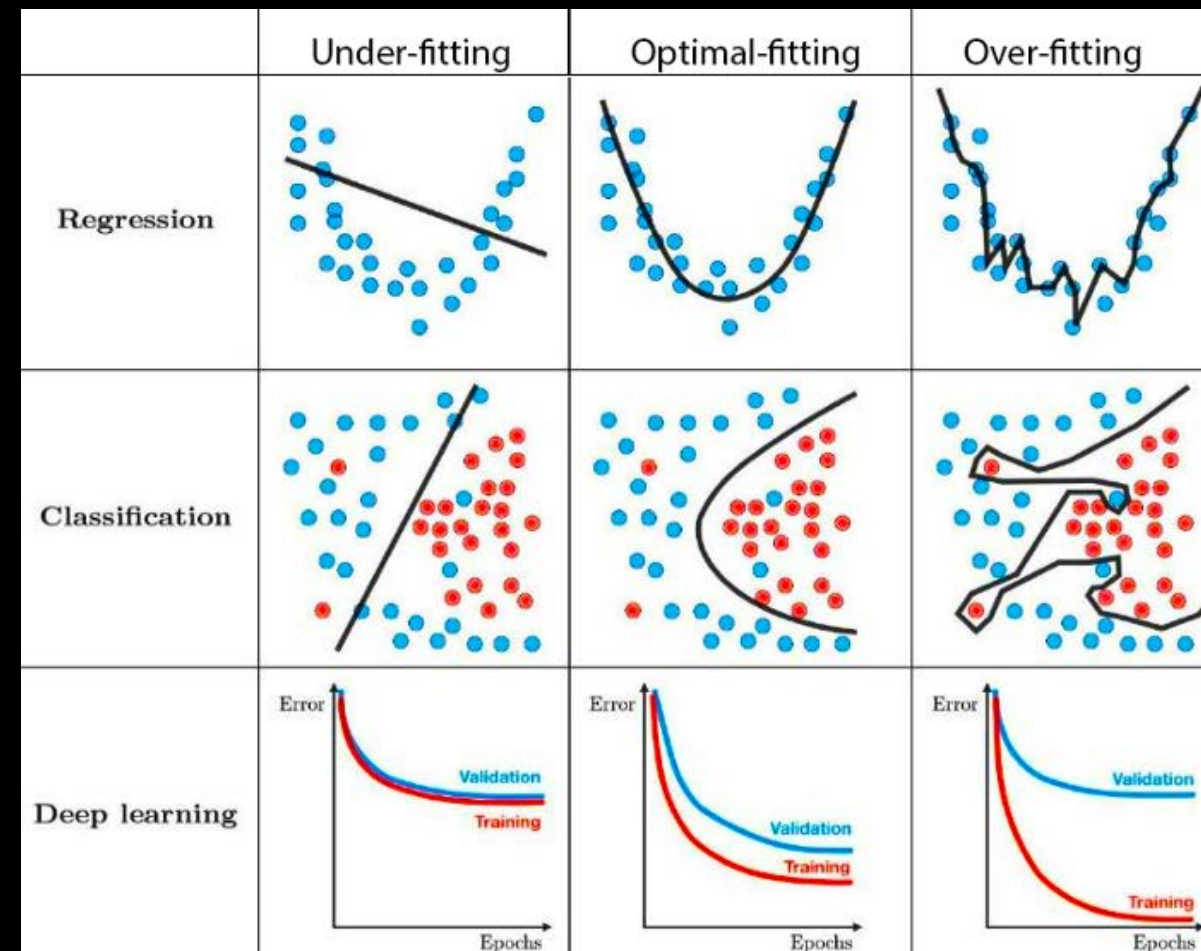


Model Complexity

Complexity & Size of data

■ Complexity ↑ for...

1. A model with more parameters
 - Degree of Freedom (DoF)
2. A model whose parameters can take a wider range of values
 - When weights can take a wider range of values
3. A model that takes more training iterations




How to deal with Underfitting/Overfitting?

Complexity & Size of data

■ Complexity ↑ for...

1. A model with more parameters
 - Degree of Freedom (DoF)
2. A model whose parameters can take a wider range of values
 - When weights can take a wider range of values
3. A model that takes more training iterations

Image credit: <https://srdas.github.io/DLBook2/ImprovingModelGeneralization.html>



1. Add more parameters
 - More hidden layers
 - More number of units per layer
3. Train longer

1. Add more data
 - ~Data augmentation
2. Weight decay, Dropout
3. Early stopping
 - **Regularization techniques (1-3)**

+Different optimization algorithms, network architectures

More data never hurt!

- In theory
 - Fewer samples \rightarrow more overfitting
 - # training data $\uparrow \rightarrow$ generalization error \downarrow
 - Absent sufficient data, simpler models may be more difficult to beat.
- But in practice...
 - Costly, Time consuming
 - e.g.) Adding one MRI brain image of a dementia patient ~\$200

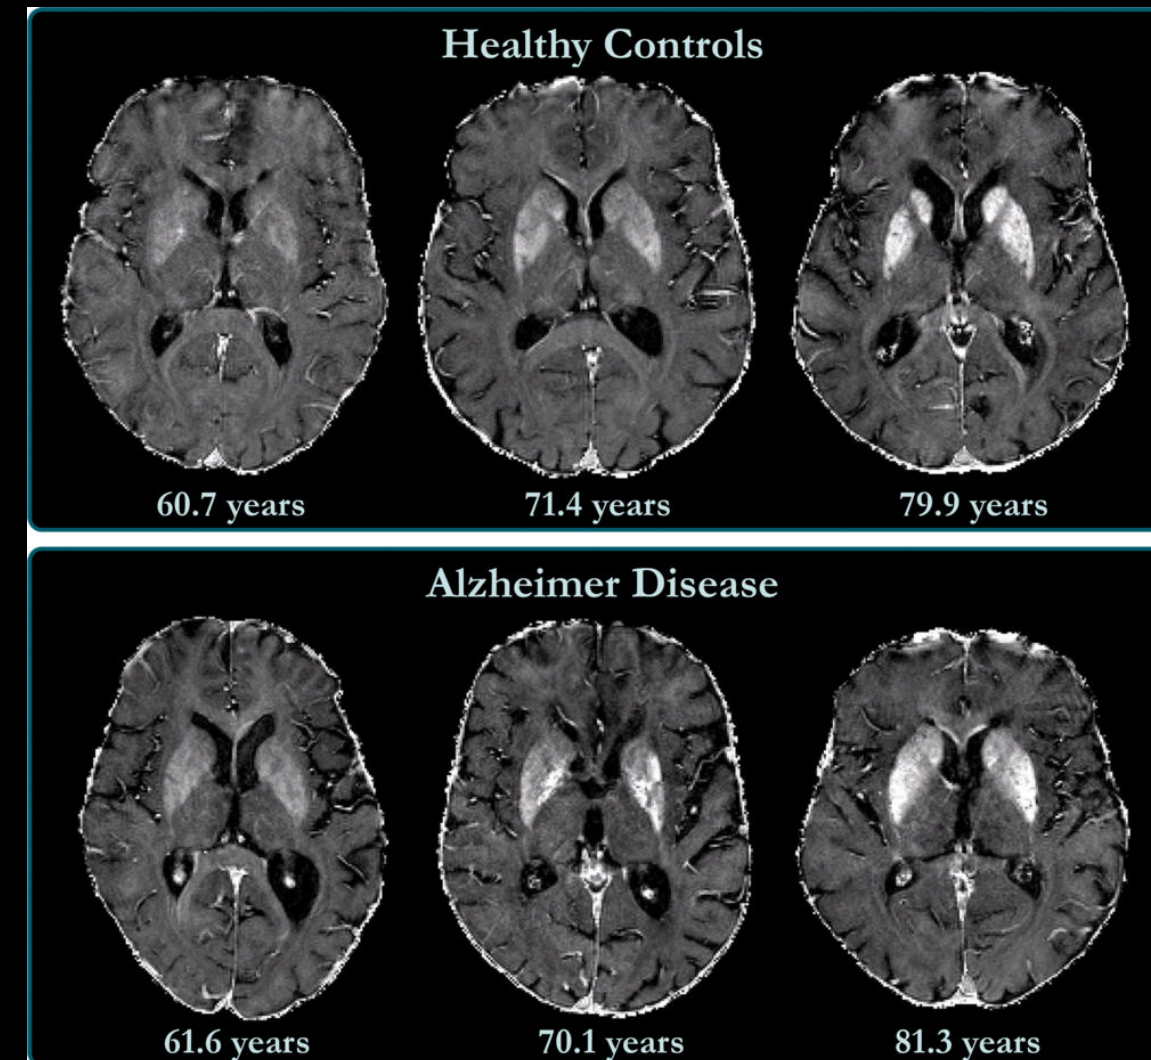



Image credit: <https://www.appliedradiology.com/communities/MR-Community/mri-shows-brain-iron-accumulation-linked-to-cognitive-deterioration-in-alzheimer-s-patients>


Data augmentation

- Techniques used to increase the amount of data
 1. Slightly modified copies of already existing data
 2. Newly created synthetic data from existing data.
- Acts as a regularizer and helps reduce overfitting



Weight Decay

- Weight decay
(L_2 regularization)
 - The most widely-used technique for regularizing parametric machine learning models.
- Among all functions f , the function $f = 0$ (assigning the value 0 to all inputs) is in some sense the **simplest**
 - → Measure the complexity of a function by its distance from 0




Weight Decay

■ Norms (Section 2.3.10)

- Norms : $\|x\|_p = (\sum_{i=1}^n |x_i|^p)^{\frac{1}{p}}$
- L1 Norms $\|x\|_1 = \sum_{i=1}^n |x_i|$
- L2 Norms $\|x\|_2 = \sqrt{\sum_{i=1}^n x_i^2}$
 - Frobenius norm (For a matrix X)
 - $\|x\|_F = \sqrt{\sum_{i=1}^m \sum_{j=1}^n x_{ij}^2}$

- Add its norm as a **penalty term** to the problem of minimizing the loss.
 - Minimizing **the prediction loss** on training sets
→ Minimizing both the **prediction loss** & **penalty term**.



Weight Decay

WeiLeong's notebook

Gradient Descent

Therefore it is natural to define a **loss function** in linear regression, as

$$\ell^i = \frac{1}{2} \left(y^{(i)} - \hat{y}^{(i)} \right)^2 \quad (6)$$

for single sample set. For entire set,

$$L(\mathbf{w}, \mathbf{b}) = \frac{1}{n} \sum_{i=1}^n \ell^i = \frac{1}{n} \sum_{i=1}^n \frac{1}{2} \left(\mathbf{y}^{(i)} - \mathbf{w}^\top \mathbf{x}^{(i)} - \mathbf{b} \right)^2 \quad (7)$$

and the goal is find

$$(\mathbf{w}^*, \mathbf{b}^*) = \underset{\mathbf{w}, \mathbf{b}}{\operatorname{argmin}} L(\mathbf{w}, \mathbf{b})$$

Besides that, we also need to control how far each update goes. Sometimes the gradient may be too extreme and the optimum value just passes off, so we need to adjust **learning rate** η , which controls the step for update process.

$$\begin{aligned} (\mathbf{w}, \mathbf{b}) &\leftarrow (\mathbf{w}, \mathbf{b}) - \eta \frac{\partial L(\mathbf{w}, \mathbf{b})}{\partial (\mathbf{w}, \mathbf{b})} \\ w_j &\leftarrow w_j - \eta \frac{1}{n} \sum_{i=1}^n \frac{\partial \ell^{(i)}(\mathbf{x}^{(i)}, y^{(i)}, \mathbf{w})}{\partial w_j} \end{aligned} \quad (8)$$

Gradient Descent with L2 Regularization

λ : Regularization constant
(Degree of regularization)


$$\begin{aligned} L'(\mathbf{w}) &= L(\mathbf{w}) + \frac{1}{2} \lambda \|\mathbf{w}\|^2 \\ &= L(\mathbf{w}) + \frac{1}{2} \lambda \sum_j w_j^2 \end{aligned}$$

$$\begin{aligned} w_j &\leftarrow w_j - \eta \frac{\partial L'(\mathbf{w})}{\partial w_j} \\ &= w_j - \eta \left(\frac{\partial L(\mathbf{w})}{\partial w_j} + \lambda w_j \right) \\ &= (1 - \eta \lambda) w_j - \eta \frac{\partial L(\mathbf{w})}{\partial w_j} \end{aligned}$$


Dropout

- Background from Bishop 1995
- Developed by Srivastava et al., 2014
 - Inject noise into each layer of the network before calculating the subsequent layer during training.
 - “Dropout” - We literally drop out some neurons during training.

MLP with one hidden layer




Hidden layer after dropout




Dropout

- During training, **randomly drop** hidden unit on some probability p from the neural network on each training iteration.
 - Activation h is replaced by h'
 - $$h' = \begin{cases} 0 & \text{Dropout} \\ h / (1 - p) & \text{Debiases each layer} \\ & \rightarrow \text{Expectation does not change} \end{cases}$$
- Can apply higher dropout probability to layers with more hidden units.
- For test / validation sets \rightarrow turn off dropout
 - We don't want to add noise for evaluation

MLP with one hidden layer



Hidden layer after dropout



Dropout

- Dropout breaks co-adaptation among neurons
 - Some neurons are highly dependent on others
- Model weights are more motivated to spread out across many hidden units
 - Not depending too much on a small number of potentially spurious associations.
 - → Similar to the effect of applying L2 regularization.




Figure 4: Test error for different architectures with and without dropout. The networks have 2 to 4 hidden layers each with 1024 to 2048 units.


How to deal with Underfitting/Overfitting?

Complexity & Size of data

■ Complexity ↑ for...

1. A model with more parameters
 - Degree of Freedom (DoF)
2. A model whose parameters can take a wider range of values
 - When weights can take a wider range of values
3. A model that takes more training iterations

Image credit: <https://srdas.github.io/DLBook2/ImprovingModelGeneralization.html>



1. Add more parameters
 - More hidden layers
 - More number of units per layer
3. Train longer

1. Add more data
 - ~Data augmentation
2. Weight decay, Dropout
3. Early stopping
 - **Regularization techniques (1-3)**

+Different optimization algorithms, network architectures


Summary

Introduction

DB - Result: Not found
ML - Result: Spiral

Under/
Overfitting

Data Sets



Data set					
Training set				Test set	
Training set				Validation set	Test set
Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Finding parameters
Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	
Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	
Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	
Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	
$Error = \frac{1}{5} \sum Err_i$					Test set Final evaluation

How to Kill
Overfitting

Complexity ↓ & Size of data ↑

